

YANKI

Geo Intelligence Platform — Competitive Analysis, Product Strategy & PRD

Building a Semrush-class Geo Intelligence platform for the post-Adobe era of local and AI search

Document type	Market research, product strategy & Product Requirements Document
Working product name	Yankı (“echo” in Turkish — your brand’s echo across the map and across AI answers)
Prepared for	Founders, product managers, designers, and engineering teams
Version	1.0
Date	August 3, 2026
Status	Planning baseline — intended as the primary planning document

Executive Summary

The SEO tooling market is in the middle of its largest structural shift since Google Panda. Three forces converged between 2024 and 2026 and created a rare window for a focused challenger:

1. The incumbent got absorbed. Adobe completed its ~\$1.9B acquisition of Semrush on April 28, 2026. Semrush is being repositioned as the “brand visibility” layer inside Adobe CX Enterprise — SEO, Generative Engine Optimization (GEO), and Agentic Search Optimization (ASO) for large enterprises. History is consistent about what happens next: enterprise integration roadmaps absorb engineering capacity, per-seat pricing drifts upward, and the long tail of SMBs, agencies, and local businesses becomes a lower priority. This is the classic moment when a focused, cheaper, faster challenger takes the segment the incumbent is walking away from.

2. Search became geographic and generative at the same time. “Near me” intent keeps growing while AI assistants (Google AI Overviews / AI Mode, ChatGPT, Perplexity, Copilot, Claude) increasingly answer local questions directly — “best dentist in Kadıköy,” “emergency plumber open now in Dubai Marina.” No mainstream platform today measures a business’s visibility across **both** the map pack geo-grid **and** AI-generated local answers in a single Share of Local Voice model. The word “GEO” itself now has two meanings — *geographic* and *generative engine optimization* — and the winning product will own both.

3. Local tooling is fragmented and mispriced. The local stack is split between generalists that treat local as an add-on (Semrush Local at \$30-60/location/month) and point tools that do one thing brilliantly (Local Falcon’s geo-grids) but nothing else. Agencies today stitch together 3-6 tools per

client. Nobody offers an integrated, multi-tenant, white-label, agency-economics platform with modern AI.

The recommendation of this report: build **Yanki**, a **Geo Intelligence Operating System** targeting local-focused digital agencies (primary ICP) and mid-market multi-location brands (secondary ICP), with three durable differentiators:

1. **Unified Visibility Index** — one score combining geo-grid map rankings, organic local rankings, and AI-answer share of voice (the “two GEOs” fused).
2. **Agency-native economics** — per-location pricing that scales *down* aggressively with volume (target \$3-8/location at scale vs. Semrush’s \$30-60), true multi-tenant workspaces, and white-label client portals from day one.
3. **Automation & agent-readiness** — prescriptive playbooks (detect → diagnose → recommend → verify) and a first-class API + MCP server so AI agents can operate the platform, not just read from it.

The report that follows delivers the full competitive landscape (Section 1), feature analysis (2), segmentation and personas (3-4), ICP selection (5), product strategy (6), a complete PRD (7), feature prioritization (8), user flows (9), multi-tenant B2B SaaS architecture (10), RBAC (11), subscription design (12), differentiation strategy (13), and a 24-month roadmap (14).

Headline plan: ship an MVP in 90 days centered on geo-grid tracking + GBP integration + white-label reporting; reach competitive parity on listings/reviews/local audits by month 6; own the AI-visibility-for-local category by month 12; expand to enterprise multi-location and an agent/data platform by month 24. Target: \$1M ARR within 18 months on agency-led growth, with COGS discipline built into a credit-metered scan economy from day one.

Table of Contents

1. **Competitive Landscape** — 11 platform profiles + comparison matrix
2. **Feature Analysis** — core (industry-standard) features & differentiators
3. **Customer Segmentation** — 12 segments with workflows, KPIs, pain points
4. **Persona Analysis** — 5 deep personas with JTBD and Semrush usage patterns
5. **Ideal Customer Profile (ICP)** — ranked ICP opportunities with rationale
6. **Product Strategy** — MVP → V1 → V2 → Vision; MoSCoW classification
7. **Product Requirements Document (PRD)** — full PRD for the selected ICP
8. **Feature Prioritization** — Tables A-D with priority, value, complexity, dependencies
9. **User Flows** — 15 flows with triggers, steps, edge cases, success states (+ Mermaid)
10. **B2B SaaS Architecture** — auth, org hierarchy, seat management, migration
11. **Roles & Permissions** — complete RBAC model with permission matrices
12. **Subscription Model** — tiers, limits, quotas, credits, add-ons
13. **Differentiation Strategy** — ranked by expected business impact
14. **Product Roadmap** — 0-3 / 3-6 / 6-12 / 12-24 months
15. **References & Assumptions**

Reading Notes & Key Assumptions

- **Pricing data:** Semrush, Ahrefs, Local Falcon, and BrightLocal prices were verified against public sources in early August 2026. Prices for Moz, SE Ranking, Similarweb, Whitespark, GeoRanker, Local

Viking and the enterprise listings vendors are stated as approximate list prices as of the research date and should be re-verified before external use. All prices USD, monthly billing unless noted.

- **Terminology:** *GBP* = Google Business Profile. *Geo-grid* = a lattice of simulated search points (e.g., 7×7, 13×13) around a location, each returning local pack rankings. *SoLV* = Share of Local Voice. *GEO* is used in both senses — geographic intelligence and Generative Engine Optimization — and the report treats the fusion of the two as the core strategic thesis. *AIV* = AI Visibility (presence/citations in AI-generated answers).
 - **Diagrams:** user journeys, flows, org hierarchy, RBAC and architecture diagrams are provided in **Mermaid syntax** (as requested) so they can be rendered directly in Notion, GitHub, Obsidian, mermaid.live, or docs pipelines.
 - **Currency of claims:** competitive feature sets change monthly in this category. Every “Product X lacks Y” claim should be treated as a snapshot dated August 2026.
 - **Cost model assumption:** the dominant variable cost in geo-grid products is SERP/Places data acquisition (each grid pin × keyword = one localized SERP fetch). The pricing architecture in Section 12 is designed around this reality.
-

1. Competitive Landscape

1.1 Market structure: three rings

The market is best understood as three concentric rings. Yanki's strategy is to start in Ring 3 (where trust and workflows are won), absorb the Ring 2 jobs agencies actually bill for, and selectively pull Ring 1 capabilities in as context — never trying to out-breadth Semrush.

Ring	Description	Players	Entry price band
Ring 1 — Visibility suites	Broad SEO/marketing intelligence; local is a module or add-on	Semrush (Adobe), Ahrefs, Moz, SE Ranking, Similarweb	\$29-\$550+/mo
Ring 2 — Local SEO platforms	Purpose-built local: listings, reviews, audits, rank tracking	BrightLocal, Whitespark, Moz Local, Uberall, Yext, SOCi, Chatmeter	\$20-\$60/loc/mo (SMB); \$200+/loc/yr (enterprise)
Ring 3 — Geo-grid specialists	Map-pack rank tracking on geographic grids; the “proof layer”	Local Falcon, Local Viking, GeoRanker, Places Scout, Nightwatch	\$20-\$200/mo credit- or location-based

A fourth, adjacent ring emerged in 2024–2026: **AI-visibility trackers** (Profound, Peec AI, Otterly.ai, plus Semrush One's AI Visibility Toolkit and Ahrefs Brand Radar). None of them are local-first. This is the seam Yanki exploits.

1.2 Platform profiles

1.2.1 Semrush (now an Adobe company)

- **Core purpose:** all-in-one online visibility suite — SEO, PPC research, content, social, competitive intelligence; now repositioned around SEO + GEO + Agentic Search Optimization inside Adobe.
- **Target market:** freelancers to enterprise; center of gravity is in-house marketing teams and agencies with \$1k+/mo tool budgets; post-acquisition pull is clearly upmarket (Adobe CX Enterprise integration).
- **Pricing positioning:** premium. Pro \$139.95 / Guru \$249.95 / Business \$499.95 per month (≈17% off annual). New **Semrush One** bundles (classic SEO + AI Visibility Toolkit): Starter \$199 / Pro+ \$299 / Advanced \$549. Extra seats billed separately (≈\$45–100/user depending on tier), which materially raises real-world cost — a 3-person Guru team lands around \$410/mo before add-ons. **Semrush Local** is a per-location add-on (Base ≈\$30/loc/mo; Premium ≈\$60/loc/mo with Map Rank Tracker heatmaps and review management).
- **Main strengths:** unmatched breadth (55+ tools); huge keyword/backlink databases; strong brand and SEO-education moat; App Center ecosystem; early and aggressive AI-visibility play; enterprise credibility now amplified by Adobe distribution.
- **Weaknesses:** complexity and feature bloat (steep learning curve, low feature utilization); real cost balloons with seats and add-ons; local is a bolted-on per-location upsell with SMB-hostile economics at scale (50 locations on Local Premium ≈ \$3,000/mo); Adobe integration risk — roadmap attention shifting to enterprise CX suites; per-project/keyword limits frustrate agencies.
- **Best customer fit:** mid-market and enterprise in-house teams that need one contract covering SEO + content + competitive intel, and agencies whose clients expect “the Semrush report.”

1.2.2 Ahrefs

- **Core purpose:** SEO data powerhouse — the industry-reference backlink index, keyword research, site audit, rank tracking; expanding into AI visibility (Brand Radar) and marketing suite ambitions.
- **Target market:** SEO professionals, content-led teams, link-building agencies; technical SEO natives.
- **Pricing positioning:** premium, usage-metered. Starter \$29 / Lite \$129 / Standard \$249 / Advanced \$449 / Enterprise \$1,499 per month; extra seats \$40-80/user/mo; credit system on lower tiers is widely disliked; API access effectively Enterprise-only or a separate high-cost subscription.
- **Main strengths:** best-in-class crawler and backlink freshness (near-real-time updates); clean UX beloved by practitioners; strong brand among professionals; excellent educational content.
- **Weaknesses:** **no meaningful local offering** — no geo-grid, no GBP management, no listings, no review tools; credit anxiety on Lite/Starter; expensive multi-seat; deliberately not a local player.
- **Best customer fit:** content/publisher SEO teams and link-focused agencies where backlinks and organic content are the growth engine.

1.2.3 Similarweb

- **Core purpose:** digital market intelligence — traffic estimation, market share, audience and app intelligence from panel + clickstream data.
- **Target market:** enterprise strategy, investment, e-commerce, and marketing intelligence teams.
- **Pricing positioning:** enterprise. Self-serve web-intelligence seats from roughly \$199/mo (annual), but real deployments are custom contracts in the tens of thousands per year.
- **Main strengths:** unique traffic/market-share dataset nobody else replicates; board-level credibility; strong for TAM sizing and competitive benchmarking.
- **Weaknesses:** not an SEO workflow tool; zero local/geo-grid capability; SMB-inaccessible pricing; estimation accuracy debated at the long tail.
- **Best customer fit:** enterprises answering “how big is the market and who is winning it,” not “how do I rank in Beşiktaş.”

1.2.4 Moz (Moz Pro + Moz Local)

- **Core purpose:** approachable SEO suite (Domain Authority, keyword research, site audit) plus a listings-sync product (Moz Local).
- **Target market:** SMBs, beginners, marketing generalists; long tail of small agencies.
- **Pricing positioning:** value/mid. Moz Pro roughly Starter \$49 / Standard \$99 / Medium \$179 / Large \$299 per month; Moz Local roughly \$14-\$33/location/mo depending on tier (approximate).
- **Main strengths:** friendliest learning curve; DA is an industry lingua franca; strong educational brand (Whiteboard Friday legacy); affordable entry.
- **Weaknesses:** data freshness and index depth lag Semrush/Ahrefs; innovation pace slowed under Ziff Davis ownership; Moz Local is listings-sync only — no geo-grid, thin reviews; declining mindshare with professionals.
- **Best customer fit:** SMBs and beginner marketers who value simplicity over depth.

1.2.5 SE Ranking

- **Core purpose:** full-stack SEO platform at aggressive value pricing; strong agency features (white label, client portals) and an early, credible AI-visibility tracker.
- **Target market:** budget-conscious agencies, freelancers, SMB in-house teams — the “80% of Ahrefs at 40% of the price” buyer.
- **Pricing positioning:** value. Roughly Essential \$65 / Pro \$119 / Business \$259 per month (approximate; frequent promos); local marketing add-on priced per location at single-digit dollars.

- **Main strengths:** exceptional price/performance; white-label + lead-gen widgets built in; flexible rank-check frequency pricing; shipped AI-search tracking (ChatGPT/Perplexity/AI Overviews) into base plans early.
- **Weaknesses:** smaller data indexes; weaker brand gravity; local module is functional but not category-leading (basic grid, lighter reviews/listings than BrightLocal).
- **Best customer fit:** small agencies and freelancers who need the whole SEO workflow, white-labeled, under \$150/mo.

1.2.6 BrightLocal

- **Core purpose:** the local SEO workbench — local rank tracking (including Local Search Grid), citation tracking/building, GBP audits, listings sync, reputation management, white-label reporting.
- **Target market:** local-focused agencies and multi-location SMBs; the default “serious local” tool in the mid-market.
- **Pricing positioning:** per-location tiers — Track \$39 / Manage \$49 / Grow \$59 per month at one location, with steep volume discounts (Grow ≈ \$449/mo at 50 locations ≈ \$9/loc; ≈\$899/mo at 100). Citation Builder is pay-as-you-go from ≈\$2/citation. 14-day free trial; white-label included.
- **Main strengths:** deepest local feature coverage in one place; agency-friendly (white label, client dashboards, lead-gen widget); trusted brand in the local community (BrightLocal Academy, surveys); sane volume economics.
- **Weaknesses:** aging UX and reporting aesthetics; review management gated to the top tier only; grid tracker less loved than Local Falcon’s; limited AI beyond “AI Insights” summaries; no AI-answer visibility tracking; no general SEO depth (thin backlinks/keywords).
- **Best customer fit:** agencies managing 5–100 local clients who need tracking + listings + reviews + reports in one bill.

1.2.7 Whitespark

- **Core purpose:** local search specialists — Local Citation Finder, Local Rank Tracker (with geo-grid), GBP management service, reputation builder; strong services arm.
- **Target market:** local SEO practitioners and small agencies; buyers of done-for-you citation/GBP services.
- **Pricing positioning:** à la carte, low-ticket software (roughly \$17–\$70+/mo per tool, approximate) plus productized services; no unified suite pricing.
- **Main strengths:** deep local expertise and community trust (Local Search Ranking Factors survey heritage); high-quality citation research; credible services.
- **Weaknesses:** fragmented product line (separate tools, separate bills); dated interfaces; limited automation, AI, and enterprise features; small team constrains velocity.
- **Best customer fit:** practitioners who want expert-grade point tools or to outsource citations/GBP work.

1.2.8 GeoRanker

- **Core purpose:** location-precise SERP data — heatmaps, local rank tracking, and above all **data/API services** (SERP retrieval by city/coordinate) for developers.
- **Target market:** developers, data teams, agencies building custom tracking.
- **Pricing positioning:** roughly \$99–\$490/mo tiers plus custom data contracts (approximate).
- **Main strengths:** raw geo-SERP data flexibility; API-first; supports many countries and custom coordinates.
- **Weaknesses:** weak end-user UX and reporting; minimal workflow features; low brand awareness; effectively a data vendor, not a platform.

- **Best customer fit:** technical buyers who want geo-SERP data to build their own layer — i.e., a potential *supplier or acqui-target* for Yankı rather than a true competitor.

1.2.9 Local Falcon

- **Core purpose:** the category-defining geo-grid rank tracker for Google Maps (and now Apple Maps): scan grids, competitor overlays, trend/campaign reports, AI analysis (Falcon AI/Assist/Agent), and proprietary metrics (SoLV, ATRP).
- **Target market:** local SEO agencies, consultants, and multi-location businesses that need visual, client-convincing proof of map visibility. Claims 100k+ organizations.
- **Pricing positioning:** pure credit model — 1 credit per grid pin. Retail packages \$24.99–\$199.99/mo (7,500–63,000 credits); enterprise \$499–\$4,999/mo (159k–1.57M credits, ~40–60% cheaper per credit); pay-as-you-go \$0.05/credit; Falcon AI analysis = 25 credits/report. Unlimited keywords/locations on all plans; API included from Basic; **already ships an MCP server** and white-label reports at enterprise tier.
- **Main strengths:** best-in-class grid UX and shareable visual reports; honest, transparent scan methodology; credit model aligns price to usage; fast AI feature velocity (AI analysis, agent, MCP); integrations (Zapier, Looker Studio).
- **Weaknesses:** **single-purpose** — no keywords research, site audit, backlinks, listings management, or review workflows; credit math confuses newcomers and punishes frequent scanning; reporting beyond grids is thin; no true multi-tenant org/roles model for large agencies; per-scan economics get expensive for dense grids × many keywords × weekly cadence.
- **Best customer fit:** anyone whose core question is “where exactly do we show up on the map, street by street” — often *alongside* another platform.

1.2.10 Local Viking

- **Core purpose:** GBP management (posting, photos, Q&A) + GeoGrid rank tracking at commodity pricing.
- **Target market:** budget agencies and freelancers managing many GBP listings.
- **Pricing positioning:** location-count tiers roughly \$20–\$149/mo covering ~1–70 GBP listings (approximate); historically undercuts everyone.
- **Main strengths:** cheap at volume; combines posting + grids; long-standing niche following.
- **Weaknesses:** dated product and slow innovation; support complaints; thin analytics/AI; limited reviews/listings depth; low ceiling for enterprise needs.
- **Best customer fit:** price-first GBP post-scheduling with basic grid proof.

1.2.11 Enterprise listings & reputation cluster (Yext, Uberall, SOCi, Chatmeter, Birdeye)

- **Core purpose:** enterprise multi-location presence — listings syndication across 100+ directories, review/reputation at scale, local pages, social; increasingly “AI-readiness” of location data.
- **Target market:** brands and franchises with 100–10,000+ locations; CMO/field-marketing buyers.
- **Pricing positioning:** opaque, contract-based; historically ~\$200–\$500+/location/year for listings, with suite bundles far above; annual contracts, sales-led.
- **Main strengths:** directory relationships and sync infrastructure; enterprise trust, SLAs, and services; deep review analytics (Chatmeter, Birdeye); SOCi’s social-at-scale.
- **Weaknesses:** expensive and slow-moving; weak practitioner-grade rank tracking (grids are shallow or absent); poor agency self-serve; innovation lag on AI-answer visibility.
- **Best customer fit:** 500+ location brands buying compliance-grade listings accuracy and reputation governance.

1.2.12 Adjacent wave: AI-visibility trackers (Profound, Peec AI, Otterly.ai, Semrush One AIV, Ahrefs Brand Radar)

Purpose-built tools that measure brand presence, citations, and sentiment inside ChatGPT, Perplexity, Gemini/AI Overviews, and Copilot answers. Pricing typically \$99-\$500+/mo (approximate). None are local-first: they track brands and prompts globally, not “how often does *this dental clinic* appear when users near *this neighborhood* ask an assistant for a dentist.” **The fusion of AI-visibility with geo-grid locality is unclaimed territory** — Yankı’s primary differentiation bet (Section 13).

1.3 Comparison matrix

Legend: ● = strong/native · ◐ = partial/basic/add-on · ○ = absent. Entry price = lowest realistic paid tier, USD/mo, Aug 2026 snapshot (◊ = approximate).

Capability	Semrush	Ahrefs	Similarweb	Moz	SE Ranking	BrightLocal	Whitespark	GeoRanker	Local Falcon	Local Viking
Entry price	\$139.95 (One \$199)	\$29-129	~\$199◊	\$49◊	\$65◊	\$39/loc	~\$17◊	~\$99◊	\$24.99	~\$
Keyword research	●	●	◐	●	●	○	○	◐	○	○
Organic rank tracking	●	●	◐	●	●	◐	◐	●	○	○
Geo-grid map tracking	◐ (Local add-on)	○	○	○	◐	●	◐	●	●●	●
Backlink analysis	●	●●	◐	●	●	○	○	○	○	○
Site audit	●	●	○	●	●	◐ (local audit)	◐	○	○	○
Listings management	◐ (Yext-powered)	○	○	● (Moz Local)	◐	●	◐	○	○	○
Review management	◐	○	○	◐	◐	● (Grow only)	◐	○	◐ (analysis)	◐
GBP posting/management	◐	○	○	○	◐	●	● (service)	○	○	●
Traffic/market intel	●	◐	●●	◐	◐	○	○	○	○	○
AI-answer visibility	● (Semrush One)	● (Brand Radar)	◐	○	●	○	○	○	◐ (AI analysis of grids)	○
AI recommendations	● (Copilot)	◐	◐	◐	◐	◐	○	○	● (Falcon AI)	○
White-label reporting	◐ (agency kit)	○	○	◐	●	●	◐	○	◐ (ent.)	◐
Multi-location mgmt	◐ (\$/loc)	○	○	◐	◐	●	◐	◐	● (unlim. locs)	●
Public API	◐ (Business+)	◐ (Ent./paid)	●	●	●	●	○	●	● (+MCP)	◐
Multi-seat RBAC	◐ (\$/seat)	◐ (\$/seat)	●	◐	●	◐	○	○	◐	○
Target center			Enterprise	SMB			Local pros	Developers		

Capability	Semrush	Ahrefs	Similarweb	Moz	SE Ranking	BrightLocal	Whitespark	GeoRanker	Local Falcon	Local Viking
	Mid-mkt/Ent	SEO pros			SMB/agency	Local agency			Local agency	Bud agency

Reading of the matrix. Two whitespace clusters jump out. (a) The **column intersection** of geo-grid ●● + AI-answer visibility ● belongs to no one. (b) The **row** “Multi-seat RBAC + white-label + multi-location at agency prices” is only partially served (BrightLocal closest, but with dated UX, gated reviews, no AIV, and shallow org model). Yankı is designed to own both.

2. Feature Analysis

2.1 Core features (industry standard)

These are table-stakes: their absence is a disqualifier in evaluations, but their presence wins no deals. Complexity: L (weeks) / M (1-2 quarters) / H (multi-quarter, specialized infra). Importance: for a **Geo Intelligence** platform specifically.

#	Feature	Purpose	User value	Business value	Complexity	Importance (Geo platform)
1	Keyword research	Discover queries, volume, difficulty, intent	Know what to target before spending effort	Drives daily engagement; data moat	H (index licensing shortcut: M)	Medium — local keyword sets are small & service-driven
2	Organic rank tracking	Monitor SERP positions over time per location/device	Proof of progress; early warning	The retention engine — daily-return habit	M	High
3	Geo-grid rank tracking	Map-pack position at every point of a lattice around a location	Street-level truth: “we’re #1 downtown, invisible 2km east”	Visual wow → demos, shares, virality; core of category	M-H (SERP fetch infra + geo simulation)	Critical — the hero feature
4	SERP feature analysis	Track packs, AI Overviews, PAA, images per query	Understand what actually occupies the results page	Differentiates data quality; feeds AIV	M	High
5	Competitor analysis	Side-by-side visibility, gaps, overlaps	“Who is beating us where, and why”	Expands seats to strategy roles; upsell	M	High
6	Site audit (local-flavored)	Crawl for technical/on-page issues incl. NAP, schema, landing pages	Actionable fix lists	Consulting-grade output justifies retainers	M	Medium-High
7	Backlink analysis	Link profiles, toxic links, gaps	Authority context for rankings	Costly to own; license (DataForSEO/ Majestic) early	H (own) / L (license)	Medium — context, not core
8	Local SEO toolkit (listings, citations, GBP)	Sync NAP data, find/fix citations, manage GBP	Consistency = ranking + trust factor	Recurring per-location revenue; sticky	M-H (directory partnerships)	Critical
9	Review management	Monitor, respond, generate reviews across sites	Reputation = conversion + ranking signal	Highest willingness-to-pay in local	M	Critical (V1)
10	Reporting (white-label)	Scheduled, branded PDF/link reports	Agencies literally resell these	Retention + acquisition (branded footer opt.)	M	Critical
11	Dashboards	Configurable at-a-glance status per client/location	One screen for Monday morning	Habit formation; exec seats	M	High
12	Alerts	Threshold/anomaly notifications (rank			L-M	High

#	Feature	Purpose	User value	Business value	Complexity	Importance (Geo platform)
		drops, review spikes, listing edits)	React in hours, not at month-end	Re-engagement channel (email/ Slack/webhook)		
13	Public API	Programmatic read/write	Custom dashboards, automation	Enterprise requirement; ecosystem lock-in	M	High
14	User & org management	Seats, roles, workspaces, SSO	Team workflows, client access control	Seat expansion = NRR engine	M	Critical (B2B core)
15	Integrations	GBP, GA4, GSC, Looker Studio, Slack, Zapier	Data where users already live	Reduces churn; partner channels	M (ongoing)	High
16	Data export	CSV/Sheets/PDF export of everything	No lock-in fear; analyst workflows	Table stakes for pro buyers	L	Medium
17	Onboarding & education	Guided setup, templates, academy	Time-to-value under 10 minutes	Activation → conversion; brand moat (cf. BrightLocal Academy)	L-M	High

2.2 Differentiating features

Where products actually win. “Build difficulty” assumes a competent 4–8 person product team.

Feature	Current owner(s)	Why it matters	Build difficulty	Yanki improvement opportunity
Geo-grid UX & shareable visual scans	Local Falcon (gold standard)	The screenshot that closes clients; visual truth beats tables	Medium	Match the polish, add before/after playback , embeddable live grids, and multi-keyword composite grids in one view
AI-answer (GEO) visibility tracking	Semrush One, Ahrefs Brand Radar, Profound/Peec (none local-first)	AI assistants now answer local intent; brands are blind to it	Medium-High (prompt sampling infra, model APIs, dedup)	Local-first AIV : sample location-conditioned prompts (“dentist near X”), extract cited businesses, compute AI-SoLV per location — fused with the map grid into one Unified Visibility Index
Share of (Local) Voice	Semrush (SoV), Local Falcon (SoLV)	Turns rankings into a market-share number executives understand	Low-Medium (on top of tracking)	Extend SoLV across pack + organic + AI answers; trend it; benchmark vs. named competitors
AI recommendations / copilot	Semrush Copilot, Falcon AI, BrightLocal AI Insights	Users don’t want data, they want “what do I do Monday”	Medium (LLM + retrieval over own data)	Prescriptive playbooks : every insight ships with an action checklist, owner, and a verification scan scheduled automatically
Predictive analytics	Nobody credible in local	Forecast rank/lead impact of actions; budget justification	High (needs data volume)	Phase 3: “expected SoLV lift if review velocity +20%” style models — defensible once data accrues
Workflow automation	Weak everywhere (Zapier glue)	Agencies run the same monthly motions per client	Medium	Native Playbooks/automations : triggers (rank drop, new review,

Feature	Current owner(s)	Why it matters	Build difficulty	Yankı improvement opportunity
				listing edit) → actions (task, draft reply, rescan, notify)
Traffic & market estimation	Similarweb, Semrush	Executive context	Very High (panel data)	Don't build; integrate/partner for context widgets
AI content generation	Semrush/ ContentShake, many	Local pages, GBP posts, review replies at scale	Low-Medium (LLM)	Ship local-aware generation (service × city pages, GBP posts, review replies in brand voice, TR/AR/DE localized) with EU AI Act Art. 50-compliant AI-content labeling metadata
White-label & client portals	SE Ranking, BrightLocal	Agencies resell the platform as their own	Medium	Full white-label (domain, email, logo) at mid-tier — not enterprise-gated; read-only client seats free
Enterprise collaboration	Similarweb, Yext	Comments, tasks, approvals across teams	Medium	Lightweight built-in tasking tied to insights (not a PM tool)
Multi-location management at scale	Yext/Uberall/SOCi (enterprise), BrightLocal (mid)	Franchises need bulk ops, roll-ups, per-location drill-down	Medium-High	Bulk import, location groups, roll-up dashboards, per-location scorecards — with self-serve pricing enterprise vendors refuse to publish
Apple Maps / multi-engine grids	Local Falcon (Apple), few others	Apple Maps + Bing + AI assistants diversify beyond Google	Medium	Fast-follow Apple Maps; add Bing Places; unify in one grid switcher
API + MCP server / agent-readiness	Local Falcon (MCP shipped), others API-only	2026 buyers automate with AI agents; agent-operable tools win integrations	Low-Medium	Full-platform MCP (read <i>and</i> safe write: run scan, draft reply, build report) + webhooks + Zapier — “the agent-native Geo platform”
Credit-metered scan economy	Local Falcon	Aligns price to COGS; enables generous “unlimited locations”	Low (billing logic)	Adopt credits for scans/AI, but bundle predictable monthly allowances into tiers to kill credit anxiety
Local pages / store locator	Uberall, Yext	Owens the landing-page side of local	High	Later (V2+): schema-perfect location pages as add-on

Synthesis. Three differentiators clear the bar of *valuable, buildable, and defensible for 12+ months*: (1) local-first AI-answer visibility fused with geo-grids (Unified Visibility Index), (2) agency-native multi-tenancy + white-label + volume economics, (3) prescriptive automation and agent-readiness. Everything else is parity work sequenced by ICP demand.

3. Customer Segmentation

Semrush's ~28M registered users span every marketing role; the paying base concentrates in far fewer segments. Below, the 12 segments that matter for a Geo Intelligence platform, benchmarked on Semrush behavior.

3.1 Segment master table

Segment	Company type / industry	Size	Tool budget / mo	Tech maturity	Decision maker	Core KPIs
S1. Local SEO agency	Digital agency, local-services niche (dental, legal, home services, med-spa, restaurants)	3-50 FTE; 10-300 client locations	\$300-\$3,000 across stack	Medium-High	Founder / Head of SEO	Client retention, rank lifts, report efficiency, margin per client
S2. Generalist digital agency	Full-service agency w/ some local clients	5-100 FTE	\$500-\$5,000	Medium	Ops/SEO lead	Utilization, upsell rate, client NPS
S3. Freelance SEO consultant	Solo/duo	1-2	\$50-\$300	Medium-High	Self	Billable hours saved, client proof, churn
S4. Multi-location brand (mid-market)	Franchise/chain: 10-250 locations (fitness, QSR, clinics, auto, retail)	50-2,000 FTE	\$1k-\$10k (local stack)	Low-Medium	Marketing Director / franchise mktg mgr	Store visits, calls, direction requests, per-location visibility, review scores
S5. Enterprise multi-location	250-10,000+ locations	2,000+ FTE	\$50k-\$500k+/yr	Medium (procurement-heavy)	VP Marketing + Procurement + IT	Brand consistency, listings accuracy, reputation, compliance
S6. SMB owner (single location)	Dentist, lawyer, salon, contractor	1-20 FTE	\$0-\$100	Low	Owner	Phone calls, bookings, "am I on the map"
S7. In-house SEO specialist	Mid-market company, any industry	100-1,000 FTE	\$200-\$1,000 (their line)	High	Specialist recommends; manager signs	Organic traffic, rankings, pipeline influence
S8. Digital Marketing Manager	SMB→mid-market generalist	20-500 FTE	\$300-\$2,000	Medium	Self w/ CMO approval	Leads, CAC, channel mix
S9. Content team	Publisher/SaaS/e-comm content org	5-50 writers	\$500-\$3,000	Medium-High	Head of Content	Traffic, topical coverage, briefs shipped
S10. E-commerce	Online retail (+ sometimes stores)	any	\$300-\$5,000	Medium-High	Growth/e-comm lead	Revenue from organic, category rankings
S11. SaaS growth team	B2B/B2C SaaS	10-500 FTE	\$500-\$5,000	High	Growth lead	Signups from organic, CAC payback

Segment	Company type / industry	Size	Tool budget / mo	Tech maturity	Decision maker	Core KPIs
S12. Resellers/ white-label platforms	Marketing platforms, hosting cos, directories reselling local SEO	any	volume contracts	High	BD/Product	Attach rate, margin, API reliability

3.2 Segment deep notes (workflow, goals, pain, success)

S1 — Local SEO agencies (the beating heart of this market). Daily workflow: morning dashboard sweep across clients → triage alerts (drops, nasty reviews, GBP suspensions) → execution blocks (GBP posts, citations, content, link outreach) → month-end reporting marathon. Goals: retain clients by *proving* visible progress; win new logos with audit wow-moments; keep tool cost per client under ~\$25. Pains: tool sprawl (BrightLocal + Local Falcon + a Semrush seat + spreadsheet glue), per-location pricing that eats margin, reporting eats 2–4 days/month, clients don’t understand rank tables (grids fix this), and now clients asking “why doesn’t ChatGPT recommend us?” with no tool to answer. Success metric: churn <3%/mo, report prep <2h/client, demonstrable rank/SoLV lift in 90 days.

S4 — Mid-market multi-location brands. Workflow: weekly visibility review by region; coordinating with franchisees/store managers on reviews and GBP hygiene; quarterly budget defense to the CFO. Goals: every location in the 3-pack for core services; review score ≥4.4 everywhere; catch listing vandalism fast. Pains: enterprise vendors (Yext/Uberall) are expensive and sales-gated; SMB tools don’t roll up; no per-location accountability view; AI assistants recommending competitors invisibly. Success: uniform per-location scorecards, exceptions surfaced automatically, spend ≤\$10/location/mo for intelligence layer.

S5 — Enterprise multi-location. Buys through 6–12 month procurement with security review (SSO, SOC 2, DPA). Cares about listings accuracy SLAs, governance, roles, audit logs, and now board-level “AI visibility” narratives. Not a launch segment — a Phase 3+ expansion once compliance posture matures.

S6 — SMB owners. Huge volume, brutal economics: low willingness to pay, high support load, high churn. Serve them *indirectly* through agencies (S1) and a generous free tier that feeds the funnel; never build the roadmap around them.

S3 — Freelancers/consultants. Semrush’s classic entry wedge. They evangelize tools inside client orgs and agencies. Serve with a cheap Starter tier + white-label lite; expect them to graduate into S1.

S7–S11 (non-local segments). These are Semrush/Ahrefs heartland. Yankı should *not* chase them in years 1–2 except where local intent exists (e-commerce with stores, SaaS with location pages). Listed here to make the strategic *exclusion* explicit.

S12 — Resellers. A Phase 3 leverage play: the same scan/AIV engine sold as API/white-label infrastructure to platforms serving SMBs (cf. GeoRanker’s business, done with better packaging).

3.3 Segment attractiveness snapshot

Segment	Market size	Willingness to pay	Acquisition cost	Fit w/ Yankı differentiation	Verdict
S1 Local agencies	Large (tens of thousands globally)	High (tool = margin)	Low-Med (community-led)	Excellent	Primary
S4 Mid-mkt multi-location	Large	High	Medium	Excellent	Secondary
S3 Freelancers	Very large	Low-Med	Low (PLG)	Good	

Segment	Market size	Willingness to pay	Acquisition cost	Fit w/ Yankı differentiation	Verdict
					Wedge/ feeder
S5 Enterprise	Medium count, huge \$	Very high	High (sales-led)	Good (later)	Phase 3+
S6 SMB direct	Massive	Low	Low but churny	Weak direct	Free-tier feeder
S12 Resellers	Medium	Medium (volume)	Medium	Good (API)	Phase 3
S2/S7-S11	Large	Med-High	High vs. incumbents	Weak (red ocean)	Deprioritize

4. Persona Analysis

Five personas cover ~90% of Yankı's early revenue. Format per the brief; "Semrush lens" answers included for each.

Persona 1 — Selin, Local SEO Lead at a 12-person agency (PRIMARY)

Attribute	Detail
Background	6 yrs SEO; started freelance, now leads local practice at an İstanbul/remote agency serving dental, legal, HVAC clients in TR + Gulf + DE
Role & responsibilities	Owns local SEO delivery for ~25 clients (~60 locations); manages 2 juniors; presells audits with founders
Daily tasks	Alert triage; GBP posting/QA; citation cleanup; grid scans before client calls; review responses; junior QA; 1-2 client calls
Main objectives	Prove impact visually; cut reporting time; keep tool spend <\$25/client; win renewals
Biggest frustrations	Runs BrightLocal and Local Falcon and a shared Semrush seat; copy-pasting between them; credit anxiety on Falcon; clients asking about ChatGPT recommendations; month-end reporting hell
Purchasing behavior	Self-serve trials; decides in <2 weeks; expenses up to ~\$300/mo alone, founder sign-off above; annual only after 3 good months
Budget ownership	Recommends & effectively controls local-stack budget (~\$800/mo)
Tool habits	Daily: dashboards, GBP, Slack. Weekly: grids, reports. Monthly: audits, citations
Reporting needs	White-label PDF + live link per client, monthly; grid images front and center; plain-language wins
Collaboration needs	Assign locations to juniors; client read-only links; comment on insights
Feature priorities	Grid tracking > white-label reports > review mgmt > alerts > AI replies > citations
JTBD	"When a client questions the retainer, help me <i>show</i> — in 30 seconds, on a map — that we moved their visibility, so the renewal is easy."

Semrush lens. Why she buys Semrush: keyword research, site audit credibility, client familiarity with the brand. Actually uses: Position Tracking, Keyword Magic, Site Audit, occasionally Listing Management. Rarely uses: Social toolkit, Advertising research, Market Explorer, most of the 55 tools. Highest perceived value: the *brand* on the report and Keyword Magic. Her Semrush seat is the first cost she'd cut if Yankı covers keywords "well enough."

Persona 2 — Mark, Field Marketing Manager, 85-location fitness franchise (SECONDARY)

Attribute	Detail
Background	9 yrs retail/franchise marketing; owns "local presence" across US + expanding to UAE/KSA
Role & responsibilities	Listings accuracy, review program, local ads coordination, franchisee enablement; reports to VP Marketing
Daily tasks	Exception review (which gyms slipped), review-response escalations, franchisee nudges, vendor wrangling

Attribute	Detail
Objectives	Every location top-3 for “gym near me”; 4.5★ network average; defend budget with location-level ROI
Frustrations	Yext renewal shock; no street-level rank truth per gym; franchisees ignore dashboards; corporate blind to AI assistants recommending rival chains
Purchasing behavior	1–3 month eval, pilot on 10 locations, then annual PO; needs SSO, security questionnaire, references
Budget ownership	Owns local-tools line (\$60–120k/yr)
Tool habits	Weekly deep-dive; monthly exec roll-up; lives in email/Teams
Reporting needs	Roll-up scorecard + per-location drill-down; exec-ready PDF; anomaly narratives auto-written
Collaboration	Regional managers get their regions only; franchisees get single-location view
Feature priorities	Multi-location roll-ups > alerts > review mgmt > grids > listings sync > SSO
JTBD	“When 85 locations compete block-by-block, show me only the exceptions and what to do, so a two-person team can manage a national footprint.”

Semrush lens. Buys Semrush Guru + Local add-on trial; chokes at \$60/location Premium × 85 = \$5,100/mo → churns to a patchwork. Uses Position Tracking & Local; never touches content/social/PPC toolkits. Perceived value ceiling: brand credibility with his VP. This persona is *structurally* mispriced by Semrush — Yanki’s pricing wedge (Section 12) exists because of Mark.

Persona 3 — Deniz, Freelance SEO consultant (WEDGE)

Background: 3 yrs in-house → freelance; 6–10 small clients across TR/EU; price-sensitive, tool-savvy, active on X/LinkedIn building in public. Objectives: look bigger than one person; automate everything; keep stack <\$150/mo. Frustrations: every tool’s “agency” features locked behind \$200+ tiers; juggling 5 logins. Purchasing: instant self-serve, monthly only, churns fast if value dips. JTBD: “Make one person deliver like an agency.” Semrush lens: uses free tier + occasional 7-day trials strategically; buys only when a client covers it; highest perceived value = Keyword Magic + the PDF logo. Deniz’s public posts are Yanki’s cheapest distribution — treat this persona as a growth channel with a product attached.

Persona 4 — Aisha, Digital Marketing Manager at a 6-clinic dental group (SMB-plus)

Background: generalist marketer, first marketing hire; low SEO depth, high accountability for bookings. Objectives: more calls/bookings per clinic; understand agency reports she currently rubber-stamps. Frustrations: tools feel built for experts; jargon; fear of being fooled by her agency. Purchasing: needs founder approval >\$100/mo; picks whatever the agency recommends *or* the tool with the clearest trial “aha.” JTBD: “Tell me in plain language whether each clinic is winning its neighborhood and what to ask my agency for.” Semrush lens: bounced off trial complexity in week 1 — the classic activation failure Yanki’s guided onboarding (Section 9, Flow F1) is designed to convert. Uses: nothing consistently. Perceived value: would pay for *clarity*, not features.

Persona 5 — Tolga, Head of SEO at a 60-person full-service agency (EXPANSION)

Background: manages 6 SEOs across national + local clients; Semrush Business + Ahrefs Standard shop. Objectives: margin per account, junior leverage, defensible retainers. Frustrations: Semrush seat costs at team scale; local clients unprofitable on the enterprise stack; wants one local sub-platform he can standardize. Purchasing: quarterly vendor reviews; negotiates; wants API + SSO + volume pricing.

JTBD: “Give my local pod a standard operating platform so juniors execute a repeatable playbook without senior babysitting.” Semrush lens: keeps Semrush for national SEO regardless — Yankı’s play is *coexistence* (the local layer), not rip-and-replace. Rarely-used Semrush features at his shop: the majority; utilization audits show <20% of tools touched monthly — the perceived-value/actual-usage gap Yankı’s focused scope exploits.

4.1 Cross-persona synthesis: why people buy Semrush vs. what they use

Question	Pattern across personas
Why buy Semrush?	Brand safety (“nobody gets fired for it”), keyword database, client expectation of the logo, one-invoice convenience
Actually used	Position Tracking, Keyword Magic/Overview, Site Audit, Backlink Overview — a ~6-tool core
Rarely used	Social suite, Advertising/PLA research, Market Explorer, most App Center apps, CRM-ish features — the long tail of 55 tools
Highest perceived value	The keyword database + the brand on reports (perceived) vs. rank tracking + audits (actual daily utility)
Implication for Yankı	Compete on the <i>used core</i> for local + the <i>unserved</i> AIV/grid fusion; never on breadth. Price against the gap between perceived and used.

5. Ideal Customer Profile (ICP)

5.1 Primary ICP: the Local-Focused Digital Agency

Dimension	Specification
Industry	Digital marketing / SEO agency with a local-services client book (healthcare, legal, home services, hospitality, fitness, automotive)
Company size	3-50 employees; sweet spot 5-20
Revenue	\$250k-\$5M agency revenue
Number of locations managed	25-250 client locations (sweet spot 40-120)
Team size using the tool	2-10 seats (analysts, account managers, founders) + client viewer seats
Annual software budget	\$5k-\$40k total; \$3k-\$15k addressable by Yanki
Existing stack	BrightLocal and/or Local Falcon, a Semrush or SE Ranking subscription, GBP, GA4/GSC, Looker Studio, Slack, a PM tool, spreadsheets as glue
Buying process	Self-serve trial → 2-4 week team eval on 3-5 real clients → founder approves card → annual after proven quarter
Decision makers	Founder/Head of SEO (economic buyer), senior analyst (champion), account managers (daily users)
Business maturity	Established delivery playbook; feels margin pressure; productizing services
Growth stage	Scaling client count faster than headcount — automation-hungry

Why this ICP first (over all others):

1. **Concentration of pain.** Every Yanki differentiator maps to a bleeding agency wound: tool sprawl → one platform; reporting time → white-label automation; per-location cost → volume pricing; “why isn’t ChatGPT recommending my client” → local AIV; junior leverage → playbooks.
2. **Multiplied distribution.** One agency = 25-250 end businesses. Agencies are also the loudest community in local SEO (X, LinkedIn, LocalU, Reddit r/localseo) — product-led + community-led growth compounds. Every white-label report is a mini-billboard.
3. **Willingness to pay with healthy unit economics.** The tool is a COGS line inside a retainer; a platform that saves 10 report-hours/month pays for itself at any tier ≤\$500.
4. **Fast feedback loops.** Agencies evaluate in days, complain loudly, and churn honestly — the ideal early-stage learning customer, unlike enterprise (slow) or SMB (silent churn).
5. **Incumbent abandonment.** Post-Adobe Semrush pulls upmarket; BrightLocal is the only credible defender and carries known gaps (UX, gated reviews, no AIV, shallow org model). A window, not a wall.

5.2 Secondary ICP: Mid-Market Multi-Location Brand

Industry: franchise/chain in fitness, QSR, clinics, home services, retail. 10-250 locations, 50-2,000 FTE, \$10M-\$500M revenue. Buyer: Marketing Director / Field Marketing Manager (Persona 2). Budget: \$10k-\$120k/yr for the local intelligence layer. Buying: pilot (5-15 locations) → security/SSO check → annual PO. Enters via self-serve Business tier, converts to Enterprise plan. Prioritized second because sales

cycles are 2–6× longer and demand compliance features (SSO, SOC 2 path, DPA) that shouldn’t gate the MVP — but the *same product* serves them, so no fork is required.

5.3 Ranked ICP opportunity list

Rank	ICP	Time-to-revenue	ACV band	CAC profile	Strategic value	Priority
1	Local-focused agencies (5–20 ppl)	Weeks	\$1.2k–\$8k	Low (PLG + community)	Feedback + distribution + logo wall	Now
2	Mid-market multi-location brands	1–2 quarters	\$6k–\$60k	Medium	ACV expansion, pricing power	Now (self-serve), push at V1
3	Freelancers/consultants	Days	\$300–\$1.2k	Very low	Funnel + evangelism	Now (Starter tier)
4	Generalist agencies’ local pods	1–2 quarters	\$3k–\$20k	Medium	Land-and-expand into big shops	Phase 2
5	Enterprise multi-location (250+)	3–4 quarters	\$50k–\$300k	High	Category credibility, NRR	Phase 3
6	Resellers / API partners	2–3 quarters	volume	Medium	Infrastructure revenue	Phase 3
7	SMB direct	n/a as paid focus	<\$600	Low but churny	Free-tier feeder only	Never primary

Regional note (founder-specific). Beachhead GTM in English-first channels (global agency community) with fast-follow localization for **Turkey** (KVKK compliance, TRY pricing sensitivity → regional pricing), the **Gulf** (Arabic RTL, Google Maps dominance + strong franchise growth, higher IAP/subscription willingness), and the **EU/DACH** (GDPR, EU AI Act Article 50 labeling for AI-generated content — turn a compliance burden into a feature: Yankı auto-labels AI-assisted outputs). These map directly to the Localization requirements in the PRD (\$7.19).

6. Product Strategy

6.1 Strategic frame

Positioning statement. *For local-focused agencies and multi-location brands, Yankı is the Geo Intelligence platform that measures and grows visibility across the map pack, organic local results, and AI-generated answers — with agency-grade multi-tenancy, white-label reporting, and automation — unlike generalist suites that treat local as an overpriced add-on and point tools that only do one job.*

Strategic bets (explicit tradeoffs):

Bet	We choose	We give up	Rationale
Depth over breadth	Local/geo excellence	National-SEO parity with Semrush/Ahrefs	Can't out-breadth Adobe; can out-focus it
Fusion over point	Grid + AIV + reviews in one index	"Best single grid tool" purity	The integrated index is the moat; point features get copied
Agencies over SMBs	Multi-tenant, white-label first	Consumer-simple SMB app polish early	Distribution math (1 agency = 100 businesses)
License data, own the layer	Buy SERP/keyword/backlink data (DataForSEO-class vendors) initially	Data-moat bragging rights early	Speed; moat comes from proprietary scan + AIV corpus over time
Credits with allowances	Metered scans inside predictable tiers	Pure unlimited promises	COGS reality of geo-fetching; kill credit anxiety with bundles

6.2 MVP → V1 → V2 → Vision

MVP (months 0-3) — “The proof machine.” One killer loop: onboard business → connect GBP → run geo-grid scan → see competitor overlay → schedule recurring scans → send a white-label report → get alerted on change. Scope: auth (email+Google), organizations/workspaces/projects (thin), GBP OAuth, grid scans (Google local pack; 3×3→13×13; SoLV/avg-rank metrics), competitor discovery from scans, scheduled scans, alerting (email/Slack), report builder (PDF + live share link, white-label logo), basic keyword suggestions (licensed data), billing (Stripe, tiers + credit allowances), CSV export. Explicitly out: listings sync, review workflows, site audit, backlinks, AIV, SSO, API.

V1 (months 3-6) — “Fire your other three tools.” Local rank tracking (organic + pack, per city), review monitoring + AI reply drafts, local/GBP audit, citation tracker (monitor first, build later), AIV **beta** (location-conditioned prompt sampling across ChatGPT/Perplexity/AI Overviews; AI-SoLV per location), Unified Visibility Index v1, client viewer seats + full white-label domains, Looker Studio connector, Zapier, public read API.

V2 (months 6-12) — “The operating system.” Playbook automations (trigger→action), AI copilot over the org's own data, listings sync (via aggregator partnership), review generation campaigns, multi-location roll-ups + location groups + scorecards, Apple Maps & Bing grids, full read/write API + **MCP server**, SSO (SAML/OIDC), audit logs, TR/DE/AR localization, mobile companion app (alerts + scan viewer).

Long-term vision (12-24+ months) — “Geo Intelligence for the agentic web.” Predictive visibility modeling (action → forecast lift), AI-agent workforce (autonomous monitoring/triage with human approval gates), local pages/store-locator module, first-party data moat (largest local AIV corpus), reseller/API platform (S12), marketplace of playbooks, enterprise compliance suite (SOC 2, DPA, data residency incl. EU/TR options).

6.3 MoSCoW classification (all features)

Priority	Features
Must Have (MVP)	Auth (email, Google), org/workspace/project shell, GBP connect, geo-grid scans + scheduling, competitor overlay, SoLV/avg-rank metrics, alerts, white-label PDF/link reports, dashboards, billing + credit allowances, CSV export, onboarding wizard
Should Have (V1)	Organic local rank tracking, review monitoring + AI replies, local audit, citation monitoring, AIV beta + Unified Index, client seats, custom domains, Looker Studio, Zapier, read API, keyword research (licensed)
Nice to Have (V2)	Listings sync, review generation, playbook automations, copilot, Apple/Bing grids, roll-ups & location groups, write API + MCP, SSO, audit logs, localization TR/DE/AR, mobile app
Future	Predictive analytics, autonomous agents, local pages module, data/reseller platform, playbook marketplace, backlink module (licensed), call/lead attribution, POS/booking integrations

Key tradeoff notes. (1) *AIV in V1-beta, not MVP*: it's the headline differentiator but needs prompt-sampling infrastructure and a quality bar; shipping a flimsy version burns the story — grid-first buys time while the corpus builds. (2) *Listings sync deferred to V2*: aggregator contracts (Yext-style networks) take months and add per-location COGS; monitoring citations delivers 60% of the perceived value at 10% of the complexity. (3) *No owned keyword/backlink index before scale*: license now, revisit at \$3M+ ARR. (4) *SSO at V2, not V1*: it gates enterprise, not the primary ICP; sequence accordingly.

7. Product Requirements Document (PRD)

Product: Yankı — Geo Intelligence Platform · **Version:** 1.0 (MVP + V1 scope; V2 flagged) · **Owner:** Founder/PM · **Status:** Draft for engineering planning · **ICP:** Local-focused agencies (primary), mid-market multi-location brands (secondary)

7.1 Vision

Every business that serves a physical area deserves to know — and control — how visible it is at every point on the map and in every AI-generated answer. Yankı is the single pane of glass where agencies and multi-location teams measure that visibility, understand why it changes, and act on it, with automation doing the repetitive 80%.

7.2 Product Goals

#	Goal	Measure (12-month)
G1	Become the daily operating tool for local agency teams	≥55% weekly active orgs; median 4+ sessions/wk/active seat
G2	Collapse the local stack (replace ≥2 tools per customer)	≥60% of paying orgs report tool consolidation (survey + churn interviews)
G3	Own “AI visibility for local” as a category	≥40% of paying orgs run AIV weekly; category term association in community surveys
G4	Make reporting near-zero effort	Median report prep <15 min/client (from 2–4 h baseline)
G5	Healthy SaaS economics	Gross margin ≥75% incl. data COGS; NRR ≥110%; logo churn <2.5%/mo by month 12

7.3 Problem Statement

Local visibility work is fragmented across 3–6 tools with incompatible data models, per-location pricing designed for enterprises, and reporting that consumes days per month. Meanwhile a growing share of local discovery happens inside AI assistants where no mainstream tool measures presence at the neighborhood level. Agencies overpay, overwork, and under-prove; multi-location brands fly blind between quarterly vendor reports. The incumbent best positioned to fix this (Semrush) is integrating into Adobe’s enterprise suite, moving away from exactly these buyers.

7.4 Target Users

Primary: agency owner/Head of SEO (economic buyer), local SEO analyst (power user), account manager (report consumer/sender). Secondary: multi-location marketing manager, regional manager, franchisee (viewer). Tertiary: end-client viewer (read-only), developer (API), AI agent (MCP client, V2).

7.5 User Stories (representative; full backlog derives from Section 9 flows)

- As an **analyst**, I can run a geo-grid scan for a keyword set around a location and see per-pin rankings with competitor overlays, so I can diagnose visibility street by street. (*MVP*)

- As an **analyst**, I can schedule recurring scans (daily/weekly/biweekly/monthly) with a chosen grid size, so trend data accrues without manual work. *(MVP)*
- As an **account manager**, I can generate a white-label report (PDF + live link) combining grids, trends, reviews, and plain-language commentary, and schedule it monthly to the client, so reporting takes minutes. *(MVP report core; V1 full)*
- As an **agency owner**, I can create a workspace per client, assign seats and roles, and invite the client as a viewer, so access mirrors my business. *(MVP thin / V1 full)*
- As an **owner**, I receive an alert when any tracked location's SoLV drops >X% or a ≤ 2 ★ review lands, so we react same-day. *(MVP)*
- As an **analyst**, I can see which businesses AI assistants recommend for my keywords near my location, and my client's AI-SoLV, so I can answer "why isn't ChatGPT recommending us?" *(V1 beta)*
- As a **marketing manager (multi-location)**, I can view a roll-up scorecard of all locations with exception flags, and drill into any location, so two people can manage 85 gyms. *(V2)*
- As a **developer/agent**, I can trigger scans, fetch results, and draft review replies via API/MCP with scoped keys, so our automations operate the platform. *(read V1 / write+MCP V2)*
- As a **billing admin**, I can see credit consumption by workspace and set caps, so no client burns the month's allowance. *(MVP basic / V1 caps)*

7.6 Functional Requirements (numbered; MoSCoW tags)

FR-1 Scanning & Tracking 1.1 Geo-grid scans: sizes 3×3–15×15, radius 0.5–50 km, circular/square lattice; engines Google local pack (MVP), Google Maps, Apple Maps, Bing (V2). *(M)* 1.2 Per-pin results: rank of target + top-20 competitors, ratings, review counts, categories. *(M)* 1.3 Metrics per scan: average rank, SoLV, top-3 coverage %, visibility delta vs. previous. *(M)* 1.4 Scheduling with quiet hours & timezone awareness; pause/resume; scan credit pre-check with estimate. *(M)* 1.5 Organic local rank tracking per keyword × location × device with SERP-feature capture. *(S)* 1.6 AIV sampling: engine-conditioned prompt templates per category/geo; entity extraction of recommended businesses; AI-SoLV per location; citation/source capture where available; weekly cadence default. *(S — V1 beta)* 1.7 Unified Visibility Index: weighted composite of grid SoLV + organic visibility + AI-SoLV, configurable weights, org-level default. *(S)*

FR-2 Business & Data Onboarding 2.1 GBP OAuth: import locations, categories, hours, photos, reviews, insights; multi-account support. *(M)* 2.2 Manual/CSV bulk location import with geocoding + dedupe. *(M basic / V2 bulk 1k+)* 2.3 Competitor management: auto-suggest from scan results; manual add; per-project list. *(M)* 2.4 Keyword setup: suggestions from category + city (licensed data), volumes where available, tagging/groups. *(M basic / S full)*

FR-3 Reputation *(S)* 3.1 Review ingestion (Google via GBP; extend to Facebook/Yelp/Tripadvisor V2), sentiment + topic tagging. 3.2 AI reply drafts in brand voice + language; approval flow; direct publish to GBP. 3.3 Review alerts by rating threshold; weekly digest. 3.4 Review generation campaigns (link/QR/SMS-email templates). *(N — V2)*

FR-4 Audits & Listings 4.1 Local audit: GBP completeness, categories vs. top competitors, NAP consistency, landing-page local schema, basic on-page. *(S)* 4.2 Citation monitor across top ~40 directories per market; accuracy diffs. *(S)* 4.3 Listings sync via aggregator partner; suppression of rogue edits. *(N — V2)*

FR-5 Insights, Alerts, Automation 5.1 Alert rules: metric threshold, anomaly (z-score), event (new review, GBP change, new competitor entering top-3). Channels: email, Slack, in-app, webhook. *(M core)* 5.2 Insight feed: ranked, plain-language findings with linked evidence. *(S)* 5.3 Playbooks: trigger → conditions → actions (create task, draft reply, schedule rescan, notify, call webhook); template gallery. *(N — V2)* 5.4 Copilot Q&A over org data with citations to screens. *(N — V2)*

FR-6 Reporting & Dashboards 6.1 Dashboard per project & per workspace: KPI tiles, trend charts, latest grid, alerts. (M) 6.2 Report builder: modular sections, commentary blocks (AI-drafted, human-edited), branding (logo MVP; colors/domain/email V1), schedule + share links with optional password. (M/S) 6.3 Roll-up dashboards & location scorecards with exception filters. (N — V2) 6.4 Looker Studio connector; CSV/XLSX export everywhere. (S / M)

FR-7 Platform 7.1 Multi-tenant orgs → workspaces → projects → locations (model in §10). (M) 7.2 RBAC per §11; client viewer seats free/unlimited. (M *thin* / S *full*) 7.3 Billing: Stripe; tiers + location packs + credit allowances + PAYG top-ups; proration; dunning. (M) 7.4 Read API v1 (keys, rate limits) (S); write API + MCP server + webhooks (N — V2). 7.5 Integrations: Slack (M), Zapier (S), GA4/GSC read (S), Looker (S). 7.6 Audit log, SSO SAML/OIDC, SCIM. (N — V2/Enterprise)

7.7 Non-functional Requirements

Category	Requirement
Performance	Dashboard TTI <2.5s P75; grid render <1.5s after data; scan completion: 7×7×1kw <90s P75, queued fairly under load
Reliability	99.9% app availability; scan pipeline success ≥99% with auto-retry; zero silent scan failures (always alert on failure)
Scalability	Design targets: 10k orgs, 500k locations, 50M scan-pins/month, 1M AIV samples/month within 18 months (see §7.13)
Data freshness	Grid data timestamped; “as-of” shown everywhere; AIV samples versioned by model/engine
Accuracy honesty	Publish scan methodology; confidence notes on AIV (sampling variance disclosed) — trust is the brand
Data retention	Raw scan detail 24 months (paid) / 3 months (free); aggregates indefinitely; export before downgrade
Compatibility	Evergreen Chrome/Safari/Firefox/Edge; responsive down to 1280 desktop, 390 mobile web (read-only acceptable MVP)

7.8 Acceptance Criteria (MVP exit — must all pass)

1. New user completes signup → GBP connect → first grid scan visible in **≤10 minutes** without docs (usability-tested with 5 target-persona users).
2. Scheduled weekly scan runs 4 consecutive weeks with zero missed runs across 50 pilot projects.
3. White-label PDF generates in <60s, renders grids correctly, and is rated “client-sendable” by ≥8/10 pilot agencies.
4. Rank-drop alert fires within 15 min of scan completion when threshold breached; Slack + email delivered.
5. Billing: upgrade, downgrade, credit top-up, and cancel each work end-to-end in test + live mode; credit ledger balances to the pin.
6. Concurrent load: 100 simultaneous 7×7 scans complete ≤5 min each; no cross-tenant data leakage under automated multi-tenant test suite.
7. All destructive actions (delete project/workspace) require confirmation and are soft-deleted for 30 days.

7.9 Success Metrics

Layer	Metric	Target M3 / M6 / M12
Acquisition	Trial starts /mo	150 / 500 / 1,500
Activation	Trial→first scheduled scan ≤48h	40% / 55% / 65%

Layer	Metric	Target M3 / M6 / M12
Conversion	Trial→paid	8% / 12% / 15%
Engagement	Weekly active orgs	45% / 50% / 55%
Value proof	Reports sent /org /mo (median)	2 / 5 / 8
Revenue	MRR	\$8k / \$40k / \$150k
Retention	Logo churn /mo · NRR	5%→ / 3.5% / <2.5% · —/100%/110%
Unit economics	Gross margin (incl. SERP+AI COGS)	65% / 72% / 78%
Category	Orgs using AIV weekly	— / 20% / 40%

7.10 Analytics Requirements

Event instrumentation from day one (PostHog or Amplitude): the funnel `signup` → `org_created` → `gbp_connected` → `scan_created` → `scan_viewed` → `scan_scheduled` → `report_sent` → `alert_actioned` → `upgraded`, plus credit-consumption events with workspace/project dimensions, feature-adoption flags, and NPS/CES micro-surveys post-report-send. Dashboards: activation funnel, credit-burn vs. plan, feature adoption by persona (role captured at onboarding), cohort retention. Server-side event bus so product analytics, billing metering, and alerting share one truth. Privacy: IP-truncation, EU data-processing agreement with the analytics vendor, respect DNT.

7.11 Security Requirements

Argon2id password hashing; TOTP MFA (M) + WebAuthn (S); OAuth tokens (GBP) encrypted at rest (KMS envelope), least-scope, auto-refresh, revocation on disconnect; org-scoped row-level security enforced at the data layer, not only app layer; API keys hashed, scoped, rotatable; rate limiting + bot defense on auth; OWASP ASVS L2 as the bar; dependency and container scanning in CI; secrets in a manager, never env-committed; DDoS/WAF at edge; incident-response runbook; SOC 2 Type I program start at V2, Type II by month 18–24; GDPR + KVKK compliant processing records, DPA template, sub-processor list public; data residency option (EU) at Enterprise.

7.12 Permissions

Full model in §11. PRD-level requirements: permissions enforce at API layer; every grant/revoke audit-logged (V2 visible log); client viewers can never see credit balances, internal notes, other workspaces, or pricing; export rights separable from view rights; AI actions (auto-reply publish) require explicit role capability + per-workspace toggle.

7.13 Scalability

Horizontally scalable stateless app tier; scan pipeline as queued workers (per-engine pools) with priority classes (interactive scan > scheduled) and per-org fairness caps; SERP/AIV fetch layer abstracted behind provider interface (multi-vendor failover: e.g., DataForSEO-class + secondary) with cost tagging per request; time-series storage for rank/pin data (partitioned Postgres → columnar/ClickHouse when pin volume demands); object storage for report artifacts; read replicas for dashboards; cache layer for hot grids; multi-region readiness deferred but schema keeps `region` on org for future residency.

7.14 API Requirements

v1 REST (+ OpenAPI spec published): auth via org-scoped keys; resources: locations, keywords, scans (create/list/get), results, reviews, reports, alerts, webhooks; pagination, idempotency keys on writes;

rate limits by plan; usage metering surfaces in-app. V2: write endpoints (trigger scan, draft/publish reply with capability checks), webhooks (scan.completed, alert.fired, review.created), **MCP server** exposing the same capability set with human-approval modes for destructive/publishing actions, llms.txt + machine-readable docs (following the emerging norm — Local Falcon already publishes both; parity is table stakes for the agent era).

7.15 Integrations

Tier 1 (V1): Google Business Profile (core), Slack, Zapier, Looker Studio, GA4/GSC (read). Tier 2 (V2): listings aggregator partner, Make/n8n, Facebook & Yelp reviews, MS Teams, HubSpot (contact/report sync). Tier 3 (later): call tracking (CallRail-class), booking systems, POS footfall, agency PM tools (ClickUp/Asana task push from playbooks).

7.16 AI Capabilities

(a) Insight narration: every scan/report gets a plain-language “what changed & why it matters” paragraph, grounded in the data, temperature-low, with numbers linked to source views. (b) Review replies: brand-voice profiles per location; multilingual (TR/EN/AR/DE launch set); human-approval default, auto-publish opt-in per workspace. (c) AIV engine: prompt-template library per business category × locale; scheduled sampling across ChatGPT, Perplexity, Google AI Overviews/AI Mode (Copilot/Claude as engines mature); entity resolution against location database; variance handling via repeated sampling; model/version recorded per sample. (d) Copilot (V2): retrieval over org’s own metrics with citation-to-screen. (e) Governance: AI outputs labeled as AI-assisted (EU AI Act Art. 50-aligned disclosure metadata on generated content), per-org AI kill-switch, no cross-tenant training, vendor DPAs; AI credit metering (§12).

7.17 Reporting Requirements

Modular report composer (sections: cover, exec summary, Unified Index trend, grid gallery w/ before-after, organic movers, reviews & sentiment, AIV highlights, work log, next actions); AI-drafted commentary editable before send; branding levels (logo → full theme → custom domain + from-address); outputs: PDF, live web link (revocable, optional password, view analytics), scheduled email; localization of report strings independent of app locale (agency in EN can send TR report); template library + org-default templates.

7.18 Notifications

Channels: in-app inbox, email, Slack, mobile push (V2 app), webhook. User-level channel preferences × event type; org-level defaults; digest bundling (immediate / hourly / daily) to prevent alert fatigue; quiet hours by timezone; alert deduplication windows; every notification deep-links to evidence.

7.19 Localization

App UI: EN at MVP; TR, DE, AR (RTL) at V2; ES/FR fast-follows. Data localization: keyword/SERP data per country-locale from day one (engine `gl/hl` correctness is core, not `i18n`). Reports localizable independently (§7.17). Formats: dates, numbers, currency per locale; TRY/EUR/GBP/AED display pricing with regional price books (PPP-adjusted TR pricing recommended). Legal: KVKK (TR), GDPR (EU), consent management, EU AI Act Art. 50 disclosure toolkit for AI-generated client-facing content.

7.20 Accessibility

WCAG 2.2 AA target: full keyboard navigation, visible focus, semantic landmarks, ARIA on custom widgets; **grid heatmaps get a colorblind-safe palette + numeric overlay mode + data-table alternative** (the hero feature must not be color-only); contrast $\geq 4.5:1$; reduced-motion respect; screen-reader-tested report share pages (clients include non-expert, assistive-tech users); EN accessibility statement; European Accessibility Act awareness for EU market claims.

7.21 Mobile Requirements

MVP: responsive web, read-optimized (dashboards, grid viewer, alerts, approve/deny review replies). V2: companion app (React Native/Expo) focused on: push alerts, grid viewer with pinch-zoom, review reply approvals, client-demo mode (offline-cached last scans for sales meetings). Explicitly not mobile: report building, org admin, billing.

8. Feature Prioritization

Complexity scale: 1 (days) – 5 (multi-quarter). Value scales: 1-5. Priority: P0 (MVP) / P1 (V1) / P2 (V2) / P3 (later).

Table A — Industry-standard features (must exist; every serious competitor has them)

Feature	Priority	User value	Business value	Complexity	Dependencies
Auth (email/Google) + org shell	P0	3	5	2	—
Geo-grid rank tracking	P0	5	5	4	SERP provider, geocoding, queue infra
Scan scheduling	P0	5	5	2	Grid engine, credits
Dashboards	P0	4	4	2	Metrics store
Alerts (threshold/event)	P0	4	4	2	Scan pipeline, notif infra
White-label reports (PDF+link)	P0	5	5	3	Metrics store, render service
Billing + credit ledger	P0	3	5	3	Stripe, metering events
CSV export	P0	3	3	1	—
Organic local rank tracking	P1	4	4	3	SERP provider
Keyword research (licensed)	P1	4	3	2	Data vendor contract
Local/GBP audit	P1	4	4	3	GBP API, crawler-lite
Review monitoring	P1	5	5	3	GBP API
Citation monitoring	P1	3	3	3	Directory scrape/partners
Read API + docs	P1	3	4	2	Auth/keys, rate limiter
User mgmt + RBAC (full)	P1	4	5	3	Org model
Integrations: Slack/Zapier/ Looker	P0/P1	3	4	2	Event bus

Table B — Competitive parity features (needed to displace the current stack)

Feature	Priority	User value	Business value	Complexity	Dependencies
Competitor overlays & side-by-side grids	P0	5	4	2	Grid engine
SoLV / top-3 coverage metrics	P0	4	4	1	Grid data
Share links w/ password + analytics	P0	4	4	2	Report service
AI reply drafts (approval flow)	P1	5	4	2	Reviews, LLM, voice profiles

Feature	Priority	User value	Business value	Complexity	Dependencies
Client viewer seats (free)	P1	4	5	2	RBAC
Custom report domains/from-address	P1	4	4	2	DNS/email infra
Listings sync (aggregator)	P2	4	4	4	Partner contract, per-loc COGS model
Review generation campaigns	P2	4	4	3	Reviews module, SMS/email vendor
Apple Maps & Bing grids	P2	3	3	3	Provider coverage
GBP posting/scheduler	P2	3	3	2	GBP API write
Looker Studio connector	P1	3	3	2	API v1
SSO (SAML/OIDC) + SCIM	P2	3	4	3	Auth service
Audit logs (visible)	P2	2	4	2	Event bus

Table C — Differentiators (the reasons to switch)

Feature	Priority	User value	Business value	Complexity	Dependencies
AIV: local AI-answer tracking + AI-SoLV	P1 (beta)	5	5	4	Prompt infra, model APIs, entity resolution
Unified Visibility Index	P1	5	5	2	Grid + organic + AIV data
Before/after grid playback & embeds	P1	4	4	2	Grid history
Volume location-pack pricing (\$3-8/loc at scale)	P0 (pricing)	5	5	1	Credit COGS model
Insight feed w/ prescriptive actions	P1	4	4	3	Metrics, LLM grounding
Playbook automations (trigger→action)	P2	5	5	3	Event bus, tasks, integrations
Multi-location roll-ups + scorecards	P2	5	5	3	Location groups, metrics agg
Full-platform MCP server + write API	P2	4	4	2	API v2, capability model
Report localization (send TR/AR/DE)	P2	4	3	2	i18n of report layer
Colorblind-safe grid + table mode	P0	3	3	1	Grid renderer

Table D — Potential disruptive features (category-creating bets)

Feature	Priority	User value	Business value	Complexity	Dependencies
Predictive visibility modeling (“if review velocity +20% → SoLV +x”)	P3	5	5	5	12+ months of cross-org data
Autonomous monitoring agents (triage + propose + approve)	P3	5	5	4	Playbooks, MCP, guardrails

Feature	Priority	User value	Business value	Complexity	Dependencies
Local AIV data network (largest neighborhood-level AI-answer corpus; benchmarks & research reports)	P3	4	5	4	AIV at scale, anonymization
Playbook marketplace (agencies publish/sell SOPs)	P3	4	4	3	Playbooks, billing rails
API/white-label infrastructure for resellers (S12)	P3	3	5	3	API v2, tenancy isolation
Local pages / store locator w/ perfect schema	P3	4	4	4	Hosting/CDN, templates
Lead attribution loop (calls/bookings ↔ visibility)	P3	5	5	4	Call/booking integrations

9. User Flows

Conventions: **Goal** (user outcome) · **Trigger** · **Steps** (S=system, U=user) · **Edge cases** · **Success state**. Mermaid provided for the four highest-stakes flows.

F1. First-time onboarding (the make-or-break flow)

Goal: from signup to a rendered geo-grid in ≤ 10 minutes. **Trigger:** clicks “Start free” from site/referral/report footer. **Steps:** U signs up (Google 1-click preferred) → S creates individual org, asks role (agency / multi-location / solo — personalizes copy & defaults) → U connects GBP (OAuth) or enters business manually → S imports locations, suggests primary location → S pre-fills 5 keywords from category+city, U edits → U picks grid preset (“Neighborhood 7×7 / 2km” default) → S shows credit estimate, runs **first scan free** regardless of plan → animated grid render + auto-generated 3-bullet insight → S prompts: schedule weekly? invite teammate? → checklist (4 items) persists in sidebar. **Edge cases:** no GBP access (manual mode with limited features flagged); OAuth fails (retry + help doc + “continue without”); zero-result keyword (suggest alternates); service-area business without storefront (offer SAB grid centered on service centroid); user abandons mid-scan (email with rendered grid link — the hook). **Success:** first scan viewed + one of {schedule, invite, report} initiated. Instrumented as the activation event.

MERMAID DIAGRAM SOURCE – render at mermaid.live / Notion / GitHub

```
flowchart TD
    A[Start free] --> B[Sign up]
    B -->|Google 1-click| C[Create individual org]
    B -->|Email| C
    C --> D[Role selection\nagency / brand / solo]
    D --> E[Connect GBP?]
    E -->|OAuth ok| F[Import locations & reviews]
    E -->|Skip / fail| G[Manual business entry]
    F --> H[Keyword suggestions ready]
    G --> H
    H --> I[Grid preset + credit estimate]
    I --> J[Run first scan FREE]
    J --> K[Grid renders + AI insight]
    K --> L[Next best action]
    L --> M[Schedule weekly scan]
    L --> N[Invite teammate]
    L --> O[Generate sample report]
    M --> P[Activated ✓]
    N --> P
    O --> P
```

F2. Organization creation & individual→org migration

Goal: solo user becomes/joins a team org without losing work. **Trigger:** “Create organization” or accepting an invite. **Steps:** U names org, sets slug/logo → S converts personal workspace into org’s first workspace (data intact, user becomes Owner) → U invites members with roles → billing transfers to org (card re-confirmation). Invite-accept path: existing individual user accepting an org invite chooses per-project: transfer into org / keep personal (workspace merge wizard). **Edge cases:** email already in another org (multi-org membership supported; switcher); transferring projects with active schedules (schedules follow project; credits switch to receiving org after confirmation); domain-match suggestion

("3 teammates from acme.com are here — join instead of create?"). **Success:** ≥ 2 active seats in the org within 7 days.

F3. Workspace & project creation

Goal: agency mirrors its client book. **Trigger:** "New workspace" (per client) / "New project" (per brand/campaign inside client). **Steps:** U creates workspace "Client X" (logo, default report branding, member access list) → creates project (name, country/locale, timezone) → attaches locations (from GBP accounts or CSV) → sets default competitors & keyword groups (clone-from-template supported). **Edge cases:** workspace limit reached (upgrade prompt with per-tier table); location already tracked in another workspace (warn duplicate spend, allow with confirm); template cloning across workspaces respects branding overrides. **Success:** project has ≥ 1 location + ≥ 3 keywords + a schedule.

F4. Domain & business onboarding (deep-dive)

Goal: connect the client's digital footprint. **Trigger:** project creation wizard step or Settings→Connections. **Steps:** GBP (locations, reviews, insights) → website domain (audit crawler verification via meta-tag/DNS optional) → GA4/GSC OAuth (optional, unlocks traffic panels) → per-location confirmation of NAP canonical record (becomes the truth for citation diffing). **Edge cases:** franchise with agency-managed GBP vs. owner-managed (invite-owner email flow to grant access); mismatched NAP between GBP and site (flag immediately — first insight!); domain verification failure (limited audit mode). **Success:** ≥ 1 data source connected beyond manual entry.

F5. Geo setup (grid configuration)

Goal: a grid that reflects real service area, at sane credit cost. **Trigger:** first scan config or per-project Geo settings. **Steps:** map view centered on location → choose shape (square/circle), size (3×3–15×15), spacing (0.25–5 km) → live pin-count × keywords × frequency = monthly credit forecast shown before save → optional custom polygon (V2) → save as project default + named presets. **Edge cases:** grid crossing water/uninhabited area (offer auto-trim); dense-urban recommendation (smaller spacing) vs. rural (wider) via heuristic; forecast exceeds plan allowance (suggest cadence/keyword trims or top-up). **Success:** saved default with forecast $\leq$ allowance, or explicit top-up.

F6. Keyword setup

Goal: the right 5–50 terms per location. **Trigger:** onboarding or Keywords tab. **Steps:** category-based suggestions + "steal from competitors" (terms competitors rank for in scans) → volumes/difficulty where licensed → group by service line → mark 3–5 as "hero" (drive Unified Index weighting & report highlights). **Edge cases:** ultra-local zero-volume terms (keep, mark "no volume data — tracked anyway"); multilingual markets (per-keyword locale, e.g., TR+EN in Dubai); keyword cap by tier (soft-block + prioritize helper). **Success:** ≥ 3 hero keywords tracked per location.

F7. Competitor setup

Goal: honest benchmark set. **Trigger:** post-first-scan prompt or Competitors tab. **Steps:** S auto-lists businesses appearing most across grid pins → U pins 3–10 as named competitors → optional competitor GBP snapshot (categories, review velocity) → competitor columns appear across dashboards/reports. **Edge cases:** brand's other branches appearing as competitors (auto-detect same-brand, exclude toggle); competitor closes/rebrands (staleness detection prompts refresh). **Success:** ≥ 3 named competitors per project.

F8. Rank tracking (daily use)

Goal: state of visibility in <60 seconds. **Trigger:** login / Slack alert deep-link. **Steps:** project dashboard → Unified Index tile + delta → grid thumbnail (click = full playback slider across scan history) → movers table (best/worst keyword-pin changes) → one-click “explain this change” (AI narration with evidence links) → optional instant rescan (credit-checked). **Edge cases:** scan-in-progress states; provider outage (banner + auto-retry status, never silent); conflicting signals (pack up, organic down) — narration must address, not hide. **Success:** session views ≥ 1 grid + acts (rescan/report/task) in $\geq 20\%$ of sessions.

F9. Reporting

Goal: client-ready proof in ≤ 15 minutes. **Trigger:** monthly schedule fires or “New report.” **Steps:** pick template → S drafts commentary per section from the period’s data → U edits tone/adds work-log items → preview (exact client view) → send (email w/ PDF + live link) or schedule → link analytics (opened, time-on-report) feed back to account manager. **Edge cases:** period with negative results (template includes honest “setback + plan” block — trust play); client language $\neq$ agency language (per-report locale); huge multi-location reports (async render + notify). **Success:** report sent; client link opened (tracked); repeat schedule active.

MERMAID DIAGRAM SOURCE – render at mermaid.live / Notion / GitHub

```
graph LR
    A[Schedule fires / New report] --> B[Template + period]
    B --> C[AI drafts section commentary]
    C --> D[Human edit & work-log]
    D --> E[Preview client view]
    E --> F{Send mode}
    F -->|Email + PDF| G[Delivered]
    F -->|Live link| G
    G --> H[Open & engagement analytics]
    H --> I[Renewal conversation armed]
```

F10. Alert management

Goal: signal without fatigue. **Trigger:** Settings→Alerts or post-onboarding prompt. **Steps:** choose events (rank drop $>N$, SoLV drop $>x\%$, review $\leq 2\star$, new top-3 competitor, GBP edit detected, scan failed) → thresholds → channels per event → digest policy → test-fire button. **Edge cases:** alert storm (auto-collapse into single digest when >5 in 10 min); channel failure (fallback to email + banner); org-default vs. personal overrides precedence clearly shown. **Success:** ≥ 1 non-default alert configured; alert→action rate tracked.

MERMAID DIAGRAM SOURCE – render at mermaid.live / Notion / GitHub

```
graph TD
    A[Scan completes / event ingested] --> B{Rule match?}
    B -->|No| Z[Store only]
    B -->|Yes| C{Dedup window?}
    C -->|Duplicate| Z
    C -->|New| D[Severity & routing]
    D --> E[In-app inbox]
    D --> F[Email]
    D --> G[Slack/webhook]
    E --> H{User action}
    F --> H
    G --> H
    H -->|Open evidence| I[Deep-link to grid/review]
    H -->|Create task| J[Playbook/task created]
    H -->|Dismiss| K[Feedback tunes anomaly model]
```

F11. AI recommendations (insight → action loop)

Goal: every insight ends in a tracked action. **Trigger:** new insight in feed / report draft / copilot ask. **Steps:** insight card (finding + evidence + confidence) → suggested actions (e.g., “add ‘emergency’ service to GBP; draft replies to 3 negative reviews; rescan in 14 days”) → U accepts → tasks created (assignee, due) → verification scan auto-scheduled → outcome recorded against the insight (creating the feedback dataset that later powers prediction). **Edge cases:** low-confidence insights labeled and demoted; action requires missing permission (request-access flow); user rejects insight (reason captured). **Success:** ≥30% of surfaced insights receive an action.

F12. Team collaboration

Goal: work happens where evidence lives. **Trigger:** @mention, task assign, or client comment. **Steps:** comments on any scan/insight/report block → @mentions notify per channel prefs → lightweight tasks (title, assignee, due, linked evidence) → workspace activity feed → client-visible comment lanes separated from internal. **Edge cases:** client accidentally sees internal thread — impossible by design (separate lanes, visual distinction, confirm-on-post-to-client); mention of user without workspace access (offer invite). **Success:** multi-seat orgs show ≥5 collaboration events/wk.

F13. Subscription upgrade

Goal: friction-free expansion at the moment of need. **Trigger:** hitting a limit (locations, credits, seats, white-label) or Billing page. **Steps:** contextual paywall states exactly what’s blocked + smallest unlock (“add 10-location pack \$49” vs. full tier jump) → proration preview → confirm → capability unlocks instantly → receipt + updated forecast. **Edge cases:** downgrade below current usage (grandfather until period end, then guided trim); failed payment (grace 7 days + dunning emails + read-only before data loss, export always allowed); annual-switch incentive surfaced after 3 paid months. **Success:** limit-hit → upgrade conversion ≥12%; zero involuntary-churn surprises.

MERMAID DIAGRAM SOURCE – render at mermaid.live / Notion / GitHub

```
graph TD
    A[Limit reached] --> B[Contextual paywall\nexact blocker + options]
    B --> C{Choice}
    C --> D[Smallest unlock]
    C --> E[Tier upgrade]
    C --> F[Not now]
    D --> G[Stripe confirm]
    E --> G
    F --> H[Reminder + usage tips]
    G --> I[Instant capability unlock]
    H --> J[Updated credit forecast]
```

F14. Multi-location roll-up review (Persona 2 weekly ritual, V2)

Goal: manage 85 locations in 30 minutes. **Trigger:** Monday scorecard email → roll-up dashboard. **Steps:** network scorecard (Unified Index distribution, worst-10 exceptions, review-risk list) → filter by region/owner → drill into a location's full project view → assign fixes to regional managers → export exec summary. **Edge cases:** franchisee-owned locations with partial data access (coverage badges); new locations auto-adopt group defaults; seasonal businesses (pause tracking without losing history). **Success:** weekly scorecard opened; ≥1 exception actioned.

F15. API / agent flow (V2)

Goal: an external agent runs the monthly motion. **Trigger:** API key/MCP connection by developer or AI agent. **Steps:** scoped key issued (read / write:scan / write:reply-draft) → agent lists projects → triggers scans → polls/webhook on completion → drafts review replies (draft-only scope) → human approves in-app → agent compiles report via API. **Edge cases:** scope violation (403 + audit event); runaway credit burn (per-key caps + kill switch); publish actions always human-gated unless org enables auto with explicit warning. **Success:** ≥10% of Business+ orgs with active API/MCP usage by month 15.

10. B2B SaaS Architecture (Multi-Tenant)

10.1 Tenancy model & entity hierarchy

One logical database with strict org-scoped row-level security (RLS) at the data layer; `org_id` on every tenant-owned row; enterprise single-tenant option deferred. Hierarchy:

Organization (billing boundary, security boundary) → **Workspaces** (client for agencies / brand-region for multi-location; branding + membership boundary) → **Projects** (a tracked business context: locale, keywords, competitors, schedules) → **Locations** (physical points; the unit of location-pack pricing) → **Data** (scans, reviews, insights, reports). **Teams** are org-level member groups grantable to workspaces in bulk (e.g., “Local Pod A”); **Departments** are an enterprise labeling layer over teams (V2, reporting only).

MERMAID DIAGRAM SOURCE – render at mermaid.live / [Notion](#) / [GitHub](#)

```
graph TD
    O[Organization\nbilling · security · credits] --> W1[Workspace: Client Acme\nbranding · membership]
    O --> W2[Workspace: Client Bolt]
    O --> T[Teams\nLocal Pod A / B]
    T -- grants --> W1
    T -- grants --> W2
    W1 --> P1[Project: Acme Dental TR]
    W1 --> P2[Project: Acme Dental UAE]
    P1 --> L1[Location: Kadıköy]
    P1 --> L2[Location: Ataşehir]
    L1 --> D1[Scans · Reviews · Insights · Reports]
```

10.2 Authentication

Capability	Spec	Phase
Signup/Login	Email+password (Argon2id), email verification, breach-password check	MVP
Google Login	OIDC; account-linking by verified email	MVP
Microsoft Login	OIDC (Entra ID personal+work)	V1
Password reset	Tokenized, 30-min expiry, single-use, session invalidation on change	MVP
MFA	TOTP (enrol/backup codes) MVP; WebAuthn/passkeys V1; org-enforced MFA policy V2	staged
Sessions	Short-lived JWT + rotating refresh; device list + remote revoke (V1)	MVP/V1
SSO (SAML 2.0 / OIDC)	Org-level IdP config, JIT provisioning, domain capture	V2 (Enterprise)
SCIM	User lifecycle sync	V2 (Enterprise)
Security events	Login alerts, impossible-travel flag, audit trail	V1/V2

10.3 Seat management

Seat model. Paid seats = internal roles (Owner, Admin, Billing Admin, Manager, Editor, Analyst). **Viewer and Guest (client) seats are free and unlimited** — a deliberate wedge against per-seat incumbents; viewers create expansion pressure, not cost.

Mechanism	Behavior
Allocation	Tier includes N paid seats; extra seats purchasable à la carte; seats are org-level, access scoped per workspace
Invitations	Email invite w/ role + workspace scope; expiring links; domain auto-join (opt-in: anyone @agency.com joins as Analyst pending approval)
Transfer	Seat reassignment moves license, not identity: departing user → deactivate, transfer owned artifacts wizard (schedules, reports, API keys) to successor
Removal	Deactivate (retain history, free the seat) vs. hard delete (GDPR path, 30-day soft window); active schedules/keys flagged for reassignment before removal completes
Permissions	Role defines capability ceiling; workspace grants define scope (details §11)
Limits	Per-tier paid-seat caps; approaching-limit banners at 80%; Owner sees seat-utilization report (last-active per seat → prune candidates)
Ownership succession	Exactly one Owner required; Owner transfer needs password + email confirm; break-glass recovery via billing-email verification

10.4 Individual accounts → organization migration

Every signup creates a **personal org** (invisible scaffolding: “Deniz’s workspace”). Migration paths: (a) *Convert*: personal org upgrades in place into a company org — data, schedules, and billing history intact; user becomes Owner. (b) *Join & merge*: accepting an invite, user picks per-project whether to transfer into the company org (credits/schedules re-bind to receiving org; audit event recorded) or keep personal. (c) *Contractor mode*: user belongs to multiple orgs simultaneously with an org switcher; identity is global, membership is per-org. Design rule: **no data hostage-taking** — export works in every state, including suspended.

10.5 System architecture (reference)

MERMAID DIAGRAM SOURCE – [render at mermaid.live](https://mermaid.live) / [Notion](#) / [GitHub](#)

```
graph LR
  subgraph Clients
    WEB[Web app SPA]
    MOB[Mobile app V2]
    API[API / MCP clients]
  end
  WEB --> GW[API Gateway\nauthZ · rate limits]
  MOB --> GW
  API --> GW
  GW --> CORE[Core app services\norgs · projects · RBAC]
  GW --> BILL[Billing svc\nStripe · credit ledger]
  CORE --> Q[(Job queue)]
  Q --> SCAN[Scan workers\ngrid · organic]
  Q --> AIV[AIV samplers\nLLM engines]
  Q --> RPT[Report renderer]
  SCAN --> PROV[SERP data providers\nmulti-vendor failover]
  AIV --> LLM[Model APIs]
  SCAN --> TSDB[(Time-series store\npins · ranks)]
  AIV --> TSDB
  CORE --> PG[(Postgres RLS\ntenant data)]
  RPT --> OBJ[(Object storage\nPDFs · assets)]
  CORE --> BUS[[Event bus]]
  BUS --> NOTIF[Notifications\nemail · Slack · webhooks]
  BUS --> ANALYT[Product analytics]
  BUS --> BILLMETER[Usage metering]
  BILLMETER --> BILL
```

Operational notes: scan workers per-engine pools with priority classes (interactive > scheduled) and per-org fairness; every provider call cost-tagged → gross margin observable per org in real time; report rendering isolated (headless render service) so heavy PDFs never block app; event bus is the single source for notifications, metering, and analytics (one truth, three consumers).

11. Roles & Permissions (RBAC)

11.1 Model

Permission = Role capability n Scope grant. Roles set what a person *can do*; workspace grants set *where*. Client-facing roles (Viewer/Guest) are structurally isolated: separate comment lanes, no internal metadata, no billing/credit visibility, ever.

MERMAID DIAGRAM SOURCE – render at mermaid.live / Notion / GitHub

```
graph TD
    OO[Org Owner] --> AD[Admin]
    AD --> MG[Manager\nworkspace-scoped admin]
    MG --> ED[Editor]
    ED --> AN[Analyst]
    AN --> VW[Viewer internal]
    VW --> GU[Guest client]
    OO -.peer.-> BA[Billing Admin\nbilling-only]
```

11.2 Permission matrix

✓ = full · 🕒 = limited/scoped · ✕ = none. “Scoped” always means: only within granted workspaces.

Capability	Owner	Admin	Billing Admin	Manager	Editor	Analyst	Viewer	Guest (client)
Org settings / delete org	✓	🕒 (no delete)	✕	✕	✕	✕	✕	✕
Billing, plan, credits purchase	✓	✕	✓	✕	✕	✕	✕	✕
View credit balances/usage	✓	✓	✓	🕒 (own ws)	✕	✕	✕	✕
Invite/remove users, assign roles	✓	✓	✕	🕒 (≤Editor, own ws)	✕	✕	✕	✕
Create/delete workspaces	✓	✓	✕	✕	✕	✕	✕	✕
Workspace settings/branding	✓	✓	✕	✓ (own)	✕	✕	✕	✕
Create/edit projects, keywords, geo, competitors	✓	✓	✕	✓	✓	🕒 (edit kw/geo, not delete)	✕	✕
Run scans / consume credits	✓	✓	✕	✓	✓	✓ (per-role daily cap opt.)	✕	✕
View dashboards & results	✓	✓	🕒 (spend views)	✓	✓	✓	✓	🕒 (curated views)
Create/edit reports	✓	✓	✕	✓	✓	🕒 (draft only)	✕	✕
Send reports / manage share links	✓	✓	✕	✓	✓	✕	✕	✕

Capability	Owner	Admin	Billing Admin	Manager	Editor	Analyst	Viewer	Guest (client)
Review reply: draft	✓	✓	✗	✓	✓	✓	✗	🔒 (suggest, if enabled)
Review reply: publish	✓	✓	✗	✓	✓	✗	✗	✗
AI features usage	✓	✓	✗	✓	✓	✓	🔒 (view outputs)	🔒 (view)
AI auto-publish toggle (per ws)	✓	✓	✗	✓	✗	✗	✗	✗
Export data (CSV/ API pulls)	✓	✓	🔒 (invoices)	✓	✓	🔒 (own-project)	✗ (toggleable)	✗
API keys create/ revoke	✓	✓	✗	🔒 (ws-scoped keys)	✗	✗	✗	✗
Webhooks/ integrations config	✓	✓	✗	✓ (own ws)	✗	✗	✗	✗
Audit log view	✓	✓	🔒 (billing events)	🔒 (own ws)	✗	✗	✗	✗
Comments: internal lane	✓	✓	✗	✓	✓	✓	✓	✗
Comments: client lane	✓	✓	✗	✓	✓	🔒	✗	✓

11.3 RBAC design rules

1. **Deny-by-default**, capabilities additive; enforcement at API layer with the UI merely reflecting (never defining) rights.
2. **Separation of duties**: Billing Admin exists so finance never needs data access; Admin never needs billing.
3. **Client isolation is architectural**, not cosmetic: Guest tokens physically cannot query internal-lane comments, credit endpoints, or cross-workspace resources (RLS predicates, not UI hiding).
4. **Export as a distinct permission** — the most common enterprise ask; toggleable down to Viewer.
5. **AI actions are twice-gated**: role capability AND per-workspace policy (e.g., auto-publish replies off by default org-wide).
6. Every permission-relevant event (grant, revoke, key issue, export, publish) emits an audit record from MVP even if the viewing UI ships at V2.
7. Custom roles deferred to Enterprise tier (clone-and-edit from templates) — eight well-chosen roles cover the ICP; premature custom-RBAC is a complexity tax.

12. Subscription Model

12.1 Pricing architecture principles

1. **Three meters, honestly disclosed:** locations (capacity), scan credits (variable COGS pass-through with margin), paid seats (collaboration value). Everything else is tier-gated features.
2. **Predictability first, metering second:** generous credit *allowances inside tiers* (unlike Local Falcon's raw credit anxiety), with transparent forecasting at configuration time (Flow F5) and PAYG top-ups that never expire.
3. **The wedge:** per-location effective cost collapses with volume — \$3-8/location at 50+ locations vs. Semrush Local's \$30-60 and enterprise vendors' opaque contracts. Client viewer seats free forever.
4. **Regional price books:** PPP-adjusted pricing for TR (and staged LATAM/SEA) to match local willingness to pay without leaking into premium markets (billing-country enforcement).

12.2 Tier table (launch proposal, USD/mo, annual –20%)

	Free	Starter \$49	Pro \$129	Business \$299	Enterprise (custom, from ~\$1k)
Best for	SMB owner / evaluation	Freelancer (Deniz)	Small agency (Selin)	Agency at scale / brand (Mark pilot)	250+ locations, compliance buyers
Locations included	1	3	15	50	250+ negotiated
Extra locations	—	\$6/loc	\$5/loc·packs of 10 (\$45)	\$4/loc·packs of 25 (\$90)	\$2-3/loc
Scan credits /mo (1 credit = 1 grid pin)	500 (≈1 weekly 7×7 ×1 kw)	7,500	30,000	120,000	≥500k
Credit top-up (never expires)	—	\$6/1k	\$5/1k	\$4/1k	\$2.5-3/1k
Tracked keywords	10	100	600	2,500	custom
Paid seats incl. (extra)	1	1 (+\$15)	5 (+\$12)	12 (+\$10)	custom
Client viewer/guest seats	—	∞ free	∞ free	∞ free	∞ free
Workspaces	1	3	25	100	∞
AI credits /mo (replies, narratives, AIV samples)	20	300	1,500	6,000	custom
AIV (AI-answer tracking)	—	🕒 1 engine	✓ 3 engines	✓ all + benchmarks	✓ + API export
Reports	1 sample	10/mo, logo	∞, full white-label	∞ + custom domain/email	∞ + templates API
Reviews module	view	monitor	monitor + AI replies	+ campaigns (V2)	+ governance
Alerts/integrations	email	+Slack	+Zapier/ Looker	+webhooks, API read/write	+SSO, SCIM, audit, SLA, DPA, residency
API	—	—	read	read/write + MCP	full + volume
Support	community	email	priority email/ chat	onboarding + chat	CSM + SLA

Illustrative economics (why the wedge holds). Assume blended SERP-fetch cost $\approx$ \$1.5–2.5 per 1,000 pins via wholesale providers at volume. Pro tier: 30k credits $\approx$ \$45–75 COGS against \$129 — healthy. A 50-location agency on Business + 120k credits pays \$299 (**$\approx$ \$6/location**) where Semrush Local Premium would bill $\approx$ \$3,000 and BrightLocal Grow $\approx$ \$449 with no AIV. Enterprise credit pricing floors are set $\geq 2.5\times$ marginal data cost to protect margin. AI credits sized so AIV weekly sampling per location $\approx$ 8–12 AI credits; model-cost drift is absorbed by adjusting engine sampling depth, never by surprise bills.

12.3 Packaging rules & upgrade psychology

Feature gates trigger contextual smallest-unlock offers (Flow F13); tier jumps reserved for capability gates (white-label $\rightarrow$ Pro; API/webhooks $\rightarrow$ Business; SSO $\rightarrow$ Enterprise). Annual pushed only after month 3 of successful monthly (higher LTV honesty). Nonprofit/edu 30% off. Agency partner program: certified agencies get margin on referred Business tiers + directory listing (the BrightLocal community play, modernized). Grandfathering promise published: limits never shrink mid-cycle.

13. Differentiation Strategy — How Yankı Beats Semrush

Ranked by expected business impact (impact × feasibility × durability):

#	Differentiation	Move	Why it wins vs. Semrush (post-Adobe)	Impact
1	Geo Intelligence fusion (product)	Unified Visibility Index across map grid + organic + AI answers, local-first AIV	Semrush One tracks brand-level AIV for enterprises; nobody answers “is my dentist recommended by ChatGPT in this neighborhood.” Category-defining, 12-18 mo copy-lag for an Adobe-integrating org	★★★★★
2	Pricing (economics)	\$3-8/loc at volume, free client seats, credits with allowances, regional price books	Semrush’s per-seat + per-location model is structurally hostile to agencies/franchises; Adobe pressure pushes it further upmarket, not down	★★★★★
3	Agency-native multi-tenancy (UX/collab)	Workspaces-per-client, white-label at Pro, client lanes, seat sanity	Semrush retrofits agency features; Yankı’s org model <i>is</i> the agency’s mental model	★★★★☆
4	Automation (workflow)	Playbooks (detect→prescribe→verify), insight-to-task loop, near-zero reporting	Incumbents ship dashboards; agencies buy outcomes-per-hour. Compounds with every template	★★★★☆
5	Agent-readiness (platform)	Full MCP + write API + webhooks + llms.txt from V2	2026 buyers wire AI agents into their stack; Adobe’s agentic story is enterprise-CX-inward, not open-tool-outward. (Local Falcon proves demand; Yankı does it platform-wide)	★★★★☆
6	Data quality & honesty (trust)	Published scan methodology, confidence intervals on AIV, “as-of” everywhere, no silent failures	Trust is the local community’s currency; incumbent black-boxes are a standing grievance	★★★★☆
7	Reporting as a growth loop	Beautiful shareable grids/report links w/ tasteful “powered by” on free tiers	Every client-facing artifact recruits the next agency — Semrush can’t risk cannibalizing brand control	★★★★☆
8	UX simplicity	10-minute activation, role-aware onboarding, plain language	Semrush’s 55-tool sprawl is its own tax; focus reads as speed	★★★★☆
9	Enterprise features, sequenced	SSO/audit/SCIM/residency at V2-Enterprise without bloating core	Lets Yankı harvest mid-market brands early while staying light	★★★☆☆
10	Localization depth	TR/AR/DE product + reports, KVKK/GDPR/AI-Act toolkit	Underserved growth markets (TR, Gulf, DACH) where global tools are EN-only and full-price	★★★☆☆

What Yankı deliberately does NOT do (anti-strategy, equally important): no PPC/social/PR toolkits, no owned crawler/backlink index before \$3M ARR, no consumer SMB app, no services arm competing with agency customers, no unlimited-everything promises that break the COGS model.

14. Product Roadmap

Phase 0→1: MVP (months 0-3) — “The proof machine”

Features	Auth + org shell, GBP connect, geo-grid engine (Google pack), scheduling, competitor overlay, SoLV metrics, alerts (email/Slack), white-label PDF+link reports, dashboards, Stripe billing + credit ledger, CSV export, onboarding wizard, colorblind-safe grid modes
Risks	SERP-provider reliability/cost drift (mitigate: dual-vendor abstraction, cost tags day one); scan accuracy skepticism (publish methodology; side-by-side validation vs. Local Falcon in beta content); scope creep toward parity features (ruthless MoSCoW)
Dependencies	SERP data vendor contract(s); GBP API approval/quota; Stripe; geocoding provider
Team	2 full-stack, 1 data/pipeline eng, 1 designer (fractional), founder-PM; ~4.5 FTE
Business outcomes	40 design-partner agencies live; ≥\$8k MRR; activation ≥40%; the “grid + report in 10 minutes” demo that sells itself

Phase 2 (months 3-6) — “Fire your other three tools” (V1)

Features	Organic local tracking, review monitoring + AI replies, local/GBP audit, citation monitor, AIV beta + Unified Visibility Index , client seats + full white-label + custom domains, Looker/Zapier, read API, keyword research (licensed), MS login, WebAuthn
Risks	AIV credibility (sampling variance → publish confidence bands, weekly cadence, never single-sample claims); review-API quotas; feature-breadth diluting polish (each module ships with one persona’s ritual perfected)
Dependencies	LLM provider agreements + cost model; keyword-data license; email-sending domain infra
Team	+1 backend, +1 AI/data eng, +0.5 support/CS; ~7 FTE
Business outcomes	\$40k MRR; tool-consolidation stories (≥2 tools replaced) as case studies; AIV beta waitlist converting; churn <3.5%

Phase 3 (months 6-12) — “The operating system” (V2)

Features	Playbook automations, copilot, listings sync (aggregator partner), review campaigns, multi-location roll-ups/scorecards/groups, Apple+Bing grids, write API + MCP server + webhooks, SSO/SCIM/audit, TR/DE/AR localization + regional price books, mobile companion app, SOC 2 Type I
Risks	Aggregator contract economics (negotiate rev-share vs. flat per-loc; keep monitor-only as fallback); enterprise motion distracting PLG core (separate lightweight sales pod); automation safety (approval-gated writes, per-key caps)
Dependencies	Aggregator partnership; IdP test matrix; app-store approvals; audit firm
Team	+2 eng, +1 PM, +1 sales-assist/CS, +0.5 compliance; ~12 FTE
Business outcomes	\$150k MRR; NRR ≥110% on location/seat expansion; first 10 mid-market brand logos; “AI visibility for local” category association in community surveys

Long-term Vision (months 12-24)

Features	
----------	--

	Predictive visibility modeling; autonomous monitoring agents (approve-gated); local pages/store-locator module; AIV data network + industry benchmark reports (PR engine); reseller/API platform (S12); playbook marketplace; SOC 2 Type II; EU/TR data-residency options; lead-attribution integrations
Risks	Model-quality bar for prediction (ship as ranges w/ backtests, never point promises); platform abuse via resellers (tiered vetting); big-co response (Adobe bundling countered by speed, price, community; BrightLocal response countered by AIV + UX + agent story)
Dependencies	12+ months proprietary cross-org data; legal for marketplace/reseller terms
Team	~18-22 FTE incl. 2-person data-science pod, partnerships lead
Business outcomes	\$1M+ ARR run-rate by month 18, \$2.5M+ by 24; recognized category leadership in Geo Intelligence; defensible data moat (largest neighborhood-level AI-answer corpus); optionality: raise, profitably compound, or strategic interest from the very platforms this report analyzed

Roadmap logic in one line

Win trust with the **grid** (proof) → win wallets with **consolidation** (parity) → win the category with **AIV fusion + automation** (differentiation) → win the future with **data + agents** (moat).

15. References & Assumptions

Verified sources (accessed Aug 2026). Adobe–Semrush acquisition completion (Adobe newsroom & Semrush press releases, Business Wire, Apr 28 2026; ~\$1.9B; SEO/GEO/ASO positioning within Adobe CX Enterprise). Semrush pricing (public pricing pages & 2026 roundups: Pro \$139.95 / Guru \$249.95 / Business \$499.95; Semrush One \$199/\$299/\$549; Local base ≈\$30/loc, premium ≈\$60/loc; per-seat add-on costs). Ahrefs pricing (Starter \$29 / Lite \$129 / Standard \$249 / Advanced \$449 / Enterprise \$1,499; seat add-ons \$40–80; credit system, API gating). Local Falcon pricing & platform (credit model \$24.99–\$199.99 retail, \$499–\$4,999 enterprise, \$0.05 PAYG, Falcon AI 25 credits/report, unlimited keywords/locations, API + MCP + llms.txt). BrightLocal pricing (Track/Manage/Grow \$39/\$49/\$59 per location entry, volume curve to ≈\$9/loc at 50–100 locations, Grow-gated review management, Citation Builder from ≈\$2/citation, white-label included).

Approximate figures (training-knowledge, re-verify before external publication): Moz, SE Ranking, Similarweb, Whitespark, GeoRanker, Local Viking, Yext/Uberall/SOCi/Chatmeter price bands; AI-visibility startup pricing; SERP wholesale cost band (\$1.5–2.5/1k pins) — validate with vendor quotes.

Key assumptions. (1) AI-assistant share of local discovery continues rising through 2027 (directionally supported by Adobe’s own 269% YoY AI-traffic datapoint). (2) Google continues permitting localized SERP retrieval via established data vendors at broadly current economics. (3) GBP API access terms remain stable for management-platform use cases. (4) Post-acquisition Semrush prioritizes Adobe-suite integration over SMB/local pricing aggression for ≥12 months. (5) Agencies’ consolidation appetite (survey-consistent across local-SEO communities) persists as budgets tighten. Each assumption has a named contingency in the relevant section (multi-vendor data layer, monitor-only fallbacks, sequencing flexibility).

Document colophon. Diagrams are embedded as Mermaid source for direct rendering in Notion/ GitHub/Obsidian/mermaid.live. This document is designed to be machine-actionable: sections 7–11 can be handed to engineering as-is; section 12’s tables seed the billing configuration; section 9’s flows seed the QA test plan.

— End of report —